

# 7

## High-Quality Routines

CC2E.COM/0778

### Contents

7.1 Valid Reasons to Create a Routine

7.2 Design at the Routine Level

7.3 Good Routine Names

7.4 How Long Can a Routine Be?

7.5 How to Use Routine Parameters

7.6 Special Considerations in the Use of Functions

7.7 Macro Routines and Inline Routines

### Related Topics

Steps in routine construction: Section 9.3

Characteristics of high-quality classes: Chapter 6

General design techniques: Chapter 5

Software architecture: Section 3.5

CHAPTER 6 DESCRIBED DETAILS of creating classes. This chapter zooms in on routines, on the characteristics that make the difference between a good routine and a bad one. If you'd rather read about high-level design issues before wading into the nitty-gritty details of individual routines, be sure to read Chapter 5, "High-Level Design in Construction" first and come back to this chapter later. If you're more interested in reading about steps to create routines (and classes), Chapter 9, "The Pseudocode Programming Process" might be a better place to start.

Before jumping into the details of high-quality routines, it will be useful to nail down two basic terms. What is a "routine?" A routine is an individual method or procedure invocable for a single purpose. Examples include a function in C++, a method in Java, a function or sub procedure in Visual Basic. For some uses, macros in C and C++ can also be thought of as routines. You can apply many of the techniques for creating a high-quality routine to these variants.

What is a *high-quality* routine? That's a harder question. Perhaps the easiest answer is to show what a high-quality routine is not. Here's an example of a low-quality routine:

### CODING HORROR

#### C++ Example Of a Low-Quality Routine

```
void HandleStuff( CORP_DATA & inputRec, int crntQtr, EMP_DATA empRec, double
    & estimRevenue, double ytdRevenue, int screenX, int screenY, COLOR_TYPE &
    newColor, COLOR_TYPE & prevColor, StatusType & status, int expenseType )
{
    int i;
    for ( i = 0; i < 100; i++ ) {
        inputRec.revenue[i] = 0;
        inputRec.expense[i] = corpExpense[ crntQtr ][ i ];
    }
    UpdateCorpDatabase( empRec );
    estimRevenue = ytdRevenue * 4.0 / (double) crntQtr;
    newColor = prevColor;
    status = SUCCESS;
    if ( expenseType == 1 ) {
        for ( i = 0; i < 12; i++ )
            profit[i] = revenue[i] - expense.type1[i];
    }
    else if ( expenseType == 2 ) {
        profit[i] = revenue[i] - expense.type2[i];
    }
    else if ( expenseType == 3 )
        profit[i] = revenue[i] - expense.type3[i];
}
```

What's wrong with this routine? Here's a hint: You should be able to find at least 10 different problems with it. Once you've come up with your own list, look at the list below:

- The routine has a bad name. *HandleStuff()* tells you nothing about what the routine does.
- The routine isn't documented. (The subject of documentation extends beyond the boundaries of individual routines and is discussed in Chapter 19, "Self-Documenting Code.")
- The routine has a bad layout. The physical organization of the code on the page gives few hints about its logical organization. Layout strategies are used haphazardly, with different styles in different parts of the routine. Compare the styles where *expenseType == 2* and *expenseType == 3*. (Layout is discussed in Chapter 18, "Layout and Style.")

- 70 ● The routine's input variable, *inputRec*, is changed. If it's an input variable,
- 71 its value should not be modified. If the value of the variable is supposed to
- 72 be modified, the variable should not be called *inputRec*.
- 73 ● The routine reads and writes global variables. It reads from *corpExpense* and
- 74 writes to *profit*. It should communicate with other routines more directly
- 75 than by reading and writing global variables.
- 76 ● The routine doesn't have a single purpose. It initializes some variables,
- 77 writes to a database, does some calculations—none of which seem to be
- 78 related to each other in any way. A routine should have a single, clearly
- 79 defined purpose.
- 80 ● The routine doesn't defend itself against bad data. If *crntQtr* equals 0, then
- 81 the expression *ytdRevenue* \* 4.0 / (double) *crntQtr* causes a divide-by-zero
- 82 error.
- 83 ● The routine uses several magic numbers: 100, 4.0, 12, 2, and 3. Magic
- 84 numbers are discussed in Section 11.1, "Numbers in General."
- 85 ● The routine uses only two fields of the *CORP\_DATA* type of parameter. If
- 86 only two fields are used, the specific fields rather than the whole structured
- 87 variable should probably be passed in.
- 88 ● Some of the routine's parameters are unused. *screenX* and *screenY* are not
- 89 referenced within the routine.
- 90 ● One of the routine's parameters is mislabeled. *prevColor* is labeled as a
- 91 reference parameter (&) even though it isn't assigned a value within the
- 92 routine.
- 93 ● The routine has too many parameters. The upper limit for an understandable
- 94 number of parameters is about 7. This routine has 11. The parameters are
- 95 laid out in such an unreadable way that most people wouldn't try to examine
- 96 them closely or even count them.
- 97 ● The routine's parameters are poorly ordered and are not documented.
- 98 (Parameter ordering is discussed in this chapter. Documentation is discussed
- 99 in Chapter 20.)

100 **CROSS-REFERENCE** The

101 class is also a good contender

102 for the single greatest

103 invention in computer

104 science. For details on how to

105 use classes effectively, See

106 Chapter 6, "Working

107 Classes."

Aside from the computer itself, the routine is the single greatest invention in computer science. The routine makes programs easier to read and easier to understand than any other feature of any programming language. It's a crime to abuse this senior statesman of computer science with code like that shown in the example above.

The routine is also the greatest technique ever invented for saving space and improving performance. Imagine how much larger your code would be if you had to repeat the code for every call to a routine instead of branching to the

108 routine. Imagine how hard it would be to make performance improvements in  
 109 the same code used in a dozen places instead of making them all in one routine.  
 110 The routine makes modern programming possible.

111 “OK,” you say, “I already know that routines are great, and I program with them  
 112 all the time. This discussion seems kind of remedial, so what do you want me to  
 113 do about it?”

114 I want you to understand that there are many valid reasons to create a routine and  
 115 that there are right ways and wrong ways to go about it. As an undergraduate  
 116 computer-science student, I thought that the main reason to create a routine was  
 117 to avoid duplicate code. The introductory textbook I used said that routines were  
 118 good because the avoidance of duplication made a program easier to develop,  
 119 debug, document, and maintain. Period. Aside from syntactic details about how  
 120 to use parameters and local variables, that was the total extent of the textbook’s  
 121 description of the theory and practice of routines. It was not a good or complete  
 122 explanation. The following sections contain a much better explanation.

## 123 7.1 Valid Reasons to Create a Routine

124 Here’s a list of valid reasons to create a routine. The reasons overlap somewhat,  
 125 and they’re not intended to make an orthogonal set.

### 126 KEY POINT

#### *Reduce complexity*

127 The single most important reason to create a routine is to reduce a program’s  
 128 complexity. Create a routine to hide information so that you won’t need to think  
 129 about it. Sure, you’ll need to think about it when you write the routine. But after  
 130 it’s written, you should be able to forget the details and use the routine without  
 131 any knowledge of its internal workings. Other reasons to create routines—  
 132 minimizing code size, improving maintainability, and improving correctness—  
 133 are also good reasons, but without the abstractive power of routines, complex  
 134 programs would be impossible to manage intellectually.

135 One indication that a routine needs to be broken out of another routine is deep  
 136 nesting of an inner loop or a conditional. Reduce the containing routine’s  
 137 complexity by pulling the nested part out and putting it into its own routine.

#### *Make a section of code readable*

138 Putting a section of code into a well-named routine is one of the best ways to  
 139 document its purpose. Instead of reading a series of statements like  
 140

```
141         if ( node <> NULL ) then
142             while ( node.next <> NULL ) do
143                 node = node.next
```

```
144         leafName = node.name
145     end while
146 else
147     leafName = ""
148 end if
149 you can read a statement like
```

```
150 leafName = GetLeafName( node )
```

151 The new routine is so short that nearly all it needs for documentation is a good  
152 name. Using a routine call instead of six lines of code makes the routine that  
153 originally contained the code less complex and documents it automatically.

#### 154 ***Avoid duplicate code***

155 Undoubtedly the most popular reason for creating a routine is to avoid duplicate  
156 code. Indeed, creation of similar code in two routines implies an error in  
157 decomposition. Pull the duplicate code from both routines, put a generic version  
158 of the common code into its own routine, and then let both call the part that was  
159 put into the new routine. With code in one place, you save the space that would  
160 have been used by duplicated code. Modifications will be easier because you'll  
161 need to modify the code in only one location. The code will be more reliable  
162 because you'll have to check only one place to ensure that the code is right.  
163 Modifications will be more reliable because you'll avoid making successive and  
164 slightly different modifications under the mistaken assumption that you've made  
165 identical ones.

#### 166 ***Hide sequences***

167 It's a good idea to hide the order in which events happen to be processed. For  
168 example, if the program typically gets data from the user and then gets auxiliary  
169 data from a file, neither the routine that gets the user data nor the routine that  
170 gets the file data should depend on the other routine's being performed first. If  
171 you commonly have two lines of code that read the top of a stack and decrement  
172 a *stackTop* variable, put them into a *PopStack()* routine. Design the system so  
173 that either could be performed first, and then create a routine to hide the  
174 information about which happens to be performed first.

#### 175 ***Hide pointer operations***

176 Pointer operations tend to be hard to read and error prone. By isolating them in  
177 routines (or a class, if appropriate), you can concentrate on the intent of the  
178 operation rather than the mechanics of pointer manipulation. Also, if the  
179 operations are done in only one place, you can be more certain that the code is  
180 correct. If you find a better data type than pointers, you can change the program  
181 without traumatizing the routines that would have used the pointers.

### *Improve portability*

Use of routines isolates nonportable capabilities, explicitly identifying and isolating future portability work. Nonportable capabilities include nonstandard language features, hardware dependencies, operating-system dependencies, and so on.

### *Simplify complicated boolean tests*

Understanding complicated boolean tests in detail is rarely necessary for understanding program flow. Putting such a test into a function makes the code more readable because (1) the details of the test are out of the way and (2) a descriptive function name summarizes the purpose of the test.

Giving the test a function of its own emphasizes its significance. It encourages extra effort to make the details of the test readable inside its function. The result is that both the main flow of the code and the test itself become clearer.

### *Improve performance*

You can optimize the code in one place instead of several places. Having code in one place means that a single optimization benefits all the routines that use that routine, whether they use it directly or indirectly. Having code in one place makes it practical to recode the routine with a more efficient algorithm or in a faster, more efficient language such as assembler.

### *To ensure all routines are small?*

No. With so many good reasons for putting code into a routine, this one is unnecessary. In fact, some jobs are performed better in a single large routine. (The best length for a routine is discussed in Section 7.4, “How Long Can a Routine Be?”)

## Operations That Seem Too Simple to Put Into Routines

### KEY POINT

One of the strongest mental blocks to creating effective routines is a reluctance to create a simple routine for a simple purpose. Constructing a whole routine to contain two or three lines of code might seem like overkill. But experience shows how helpful a good small routine can be.

Small routines offer several advantages. One is that they improve readability. I once had the following single line of code in about a dozen places in a program:

### Pseudocode Example of a Calculation

```
Points = deviceUnits * ( POINTS_PER_INCH / DeviceUnitsPerInch() )
```

This is not the most complicated line of code you’ll ever read. Most people would eventually figure out that it converts a measurement in device units to a

218 measurement in points. They would see that each of the dozen lines did the same  
219 thing. It could have been clearer, however, so I created a well-named routine to  
220 do the conversion in one place:

---

### 221 Pseudocode Example of a Calculation Converted to a Function

```
222 DeviceUnitsToPoints( deviceUnits Integer ): Integer;  
223 begin  
224     DeviceUnitsToPoints = deviceUnits *  
225         ( POINTS_PER_INCH / DeviceUnitsPerInch() )  
226 end function
```

227 When the routine was substituted for the inline code, the dozen lines of code all  
228 looked more or less like this one:

---

### 229 Pseudocode Example of a Function Call to a Calculation Function

```
230 points = DeviceUnitsToPoints( deviceUnits )  
231 which was more readable—even approaching self-documenting.
```

232 This example hints at another reason to put small operations into functions:  
233 Small operations tend to turn into larger operations. I didn't know it when I  
234 wrote the routine, but under certain conditions and when certain devices were  
235 active, *DeviceUnitsPerInch()* returned 0. That meant I had to account for division  
236 by zero, which took three more lines of code:

---

### 237 Pseudocode Example of a Calculation that Expands Under Maintenance

```
238 DeviceUnitsToPoints( deviceUnits: Integer ): Integer;  
239     if ( DeviceUnitsPerInch() <> 0 )  
240         DeviceUnitsToPoints = deviceUnits *  
241             ( POINTS_PER_INCH / DeviceUnitsPerInch() )  
242     else  
243         DeviceUnitsToPoints = 0  
244     end if  
245 end function
```

246 If that original line of code had still been in a dozen places, the test would have  
247 been repeated a dozen times, for a total of 36 new lines of code. A simple routine  
248 reduced the 36 new lines to 3.

## 249 Summary of Reasons to Create a Routine

250 Here's a summary list of the valid reasons for creating a routine:

- 251 ● Reduce complexity
- 252 ● Make a section of code readable
- 253 ● Avoid duplicate code

- 254 • Hide sequences
- 255 • Hide pointer operations
- 256 • Improve portability
- 257 • Simplify complicated boolean tests
- 258 • Improve performance

259 In addition, many of the reasons to create a class are also good reasons to create  
260 a routine:

- 261 • Isolate complexity
- 262 • Hide implementation details
- 263 • Limit effects of changes
- 264 • Hide global data
- 265 • Make central points of control
- 266 • Facilitate reusable code
- 267 • To accomplish a specific refactoring

## 268 7.2 Design at the Routine Level

269 The concept of cohesion has been largely superceded by the concept of  
270 abstraction at the class level, but cohesion is still alive and well as the workhorse  
271 design heuristic at the individual-routine level.

272 **CROSS-REFERENCE** For  
273 a discussion of cohesion in  
274 general, see “Aim for Strong  
275 Cohesion” in Section 5.3.  
276 For routines, cohesion refers to how closely the operations in a routine are  
277 related. Some programmers prefer the term “strength”: How strongly related are  
278 the operations in a routine? A function like *Cosine()* is perfectly cohesive  
because the whole routine is dedicated to performing one function. A function  
like *CosineAndTan()* has lower cohesion because it tries to do more than one  
thing. The goal is to have each routine do one thing well and not do anything  
else.

279 The idea of cohesion was introduced in a paper by Wayne Stevens, Glenford  
280 Myers, and Larry Constantine (1974). Other, more modern concepts including  
281 abstraction and encapsulation tend to yield more insight at the class level, but  
282 cohesion is still a workhorse concept for the design of routines.

283 **HARD DATA**  
284 The payoff is higher reliability. One study of 450 routines found that 50 percent  
285 of the highly cohesive routines were fault free, whereas only 18 percent of  
routines with low cohesion were fault free (Card, Church, and Agresti 1986).



286 Another study of a different 450 routines (which is just an unusual coincidence)  
287 found that routines with the highest coupling-to-cohesion ratios had 7 times as  
288 many errors as those with the lowest coupling-to-cohesion ratios and were 20  
289 times as costly to fix (Selby and Basili 1991).

290 Discussions about cohesion typically refer to several levels of cohesion.  
291 Understanding the concepts is more important than remembering specific terms.  
292 Use the concepts as aids in thinking about how to make routines as cohesive as  
293 possible.

294 *Functional cohesion* is the strongest and best kind of cohesion, occurring when a  
295 routine performs one and only one operation. Examples of highly cohesive  
296 routines include *sin()*, *GetCustomerName()*, *EraseFile()*,  
297 *CalculateLoanPayment()*, and *AgeFromBirthday()*. Of course, this evaluation of  
298 their cohesion assumes that the routines do what their names say they do—if  
299 they do anything else, they are less cohesive and poorly named.

300 Several other kinds of cohesion are normally considered to be less than ideal:

301 *Sequential cohesion* exists when a routine contains operations that must be  
302 performed in a specific order, that share data from step to step, and that don't  
303 make up a complete function when done together.

304 An example of sequential cohesion is a routine that calculates an employee's age  
305 and time to retirement, given a birth date. If the routine calculates the age and  
306 then uses that result to calculate the employee's time to retirement, it has  
307 sequential cohesion. If the routine calculates the age and then calculates the time  
308 to retirement in a completely separate computation that happens to use the same  
309 birth-date data, it has only communicational cohesion.

310 How would you make the routine functionally cohesive? You'd create separate  
311 routines to compute an employee's age given a birth date, and time to retirement  
312 given a birth date. The time-to-retirement routine could call the age routine.  
313 They'd both have functional cohesion. Other routines could call either routine or  
314 both routines.

315 *Communicational cohesion* occurs when operations in a routine make use of the  
316 same data and aren't related in any other way. If a routine prints a summary  
317 report and then reinitializes the summary data passed into it, the routine has  
318 communicational cohesion; the two operations are related only by the fact that  
319 they use the same data.

320 To give this routine better cohesion, the summary data should be reinitialized  
321 close to where it's created, which shouldn't be in the report-printing routine.

322 Split the operations into individual routines. The first prints the report. The  
323 second reinitializes the data, close to the code that creates or modifies the data.  
324 Call both routines from the higher-level routine that originally called the  
325 communicationally cohesive routine.

326 *Temporal cohesion* occurs when operations are combined into a routine because  
327 they are all done at the same time. Typical examples would be *Startup()*,  
328 *CompleteNewEmployee()*, and *Shutdown()*. Some programmers consider  
329 temporal cohesion to be unacceptable because it's sometimes associated with  
330 bad programming practices such as having a hodgepodge of code in a *Startup()*  
331 routine.

332 To avoid this problem, think of temporal routines as organizers of other events.  
333 The *Startup()* routine, for example, might read a configuration file, initialize a  
334 scratch file, set up a memory manager, and show an initial screen. To make it  
335 most effective, have the temporally cohesive routine call other routines to  
336 perform specific activities rather than performing the operations directly itself.  
337 That way, it will be clear that the point of the routine is to orchestrate activities  
338 rather than to do them directly.

339 This example raises the issue of choosing a name that describes the routine at the  
340 right level of abstraction. You could decide to name the routine  
341 *ReadConfigFileInitScratchFileEtc()*, which would imply that the routine had  
342 only coincidental cohesion. If you name it *Startup()*, however, it would be clear  
343 that it had a single purpose and clear that it had functional cohesion.

344 The remaining kinds of cohesion are generally unacceptable. They result in code  
345 that's poorly organized, hard to debug, and hard to modify. If a routine has bad  
346 cohesion, it's better to put effort into a rewrite to have better cohesion than  
347 investing in a pinpoint diagnosis of the problem. Knowing what to avoid can be  
348 useful, however, so here are the unacceptable kinds of cohesion:

349 *Procedural cohesion* occurs when operations in a routine are done in a specified  
350 order. An example is a routine that gets an employee name, then an address, and  
351 then a phone number. The order of these operations is important only because it  
352 matches the order in which the user is asked for the data on the input screen.  
353 Another routine gets the rest of the employee data. The routine has procedural  
354 cohesion because it puts a set of operations in a specified order and the  
355 operations don't need to be combined for any other reason.

356 To achieve better cohesion, put the separate operations into their own routines.  
357 Make sure that the calling routine has a single, complete job:  
358 *GetEmployeeData()* rather than *GetFirstPartOfEmployeeData()*. You'll probably  
359 need to modify the routines that get the rest of the data too. It's common to

360 modify two or more original routines before you achieve functional cohesion in  
361 any of them.

362 *Logical cohesion* occurs when several operations are stuffed into the same  
363 routine and one of the operations is selected by a control flag that's passed in.  
364 It's called logical cohesion because the control flow or "logic" of the routine is  
365 the only thing that ties the operations together—they're all in a big *if* statement  
366 or *case* statement together. It isn't because the operations are logically related in  
367 any other sense. Considering that the defining attribute of logical cohesion is that  
368 the operations are unrelated, a better name might *illogical cohesion*.

369 One example would be an *InputAll()* routine that input customer names,  
370 employee time-card information, or inventory data depending on a flag passed to  
371 the routine. Other examples would be *ComputeAll()*, *EditAll()*, *PrintAll()*, and  
372 *SaveAll()*. The main problem with such routines is that you shouldn't need to  
373 pass in a flag to control another routine's processing. Instead of having a routine  
374 that does one of three distinct operations, depending on a flag passed to it, it's  
375 cleaner to have three routines, each of which does one distinct operation. If the  
376 operations use some of the same code or share data, the code should be moved  
377 into a lower-level routine and the routines should be packaged into a class.

378 **CROSS-REFERENCE** While the routine might have  
379 better cohesion, a higher-level design issue is whether  
380 the system should be using a  
381 *case* statement instead of  
382 polymorphism. For more on  
383 this issue, see "Replace  
384 conditionals with  
385 polymorphism (especially  
386 repeated *case* statements)" in  
387 Section 24.4.

It's usually all right, however, to create a logically cohesive routine if its code consists solely of a series of *if* or *case* statements and calls to other routines. In such a case, if the routine's only function is to dispatch commands and it doesn't do any of the processing itself, that's usually a good design. The technical term for this kind of routine is "event handler." An event handler is often used in interactive environments such as the Apple Macintosh and Microsoft Windows.

*Coincidental cohesion* occurs when the operations in a routine have no discernible relationship to each other. Other good names are "no cohesion" or "chaotic cohesion." The low-quality C++ routine at the beginning of this chapter had coincidental cohesion. It's hard to convert coincidental cohesion to any better kind of cohesion—you usually need to do a deeper redesign and reimplement.

None of these terms are magical or sacred. Learn the ideas rather than the terminology. It's nearly always possible to write routines with functional cohesion, so focus your attention on functional cohesion for maximum benefit.

393

## 7.3 Good Routine Names

394 **CROSS-REFERENCE** For  
395 details on naming variables,  
396 see Chapter 11, “The Power  
of Variable Names.”

397

398

399

400

401

402

403

A good name for a routine clearly describes everything the routine does. Here are guidelines for creating effective routine names.

### *Describe everything the routine does*

In the routine’s name, describe all the outputs and side effects. If a routine computes report totals and opens an output file, *ComputeReportTotals()* is not an adequate name for the routine. *ComputeReportTotalsAndOpenOutputFile()* is an adequate name but is too long and silly. If you have routines with side effects, you’ll have many long, silly names. The cure is not to use less-descriptive routine names; the cure is to program so that you cause things to happen directly rather than with side effects.

404 **CROSS-REFERENCE** For  
405 details on creating good  
406 variable names, see Chapter  
407 11, “The Power of Variable  
Names.”

408

409

410

### *Avoid meaningless or wishy-washy verbs*

Some verbs are elastic, stretched to cover just about any meaning. Routine names like *HandleCalculation()*, *PerformServices()*, *ProcessInput()*, and *DealWithOutput()* don’t tell you what the routines do. At the most, these names tell you that the routines have something to do with calculations, services, input, and output. The exception would be when the verb “handle” was used in the specific technical sense of handling an event.

### 411 **KEY POINT**

412

413

Sometimes the only problem with a routine is that its name is wishy-washy; the routine itself might actually be well designed. If *HandleOutput()* is replaced with *FormatAndPrintOutput()*, you have a pretty good idea of what the routine does.

414

415

416

417

418

In other cases, the verb is vague because the operations performed by the routine are vague. The routine suffers from a weakness of purpose, and the weak name is a symptom. If that’s the case, the best solution is to restructure the routine and any related routines so that they all have stronger purposes and stronger names that accurately describe them.

419

420

421

422

423

424

425

### *Make names of routines as long as necessary*

Research shows that the optimum average length for a variable name is 9 to 15 characters. Routines tend to be more complicated than variables, and good names for them tend to be longer. Michael Rees of the University of Southampton thinks that an average of 20 to 35 characters is a good nominal length (Rees 1982). An average length of 15 to 20 characters is probably more realistic, but clear names that happened to be longer would be fine.

426 **CROSS-REFERENCE** For  
 427 the distinction between  
 428 procedures and functions, see  
 429 Section 7.6, “Special  
 430 Considerations in the Use of  
 Functions” later in this  
 chapter.

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

450

451 **CROSS-REFERENCE** For  
 452 a similar list of opposites in  
 variable names, see  
 453 “Common Opposites in  
 Variable Names” in Section  
 454 11.1.

455

456

457

458

459

460

461

### *To name a function, use a description of the return value*

A function returns a value, and the function should be named for the value it returns. For example, *cos()*, *customerId.Next()*, *printer.IsReady()*, and *pen.CurrentColor()* are all good function names that indicate precisely what the functions return.

### *To name a procedure, use a strong verb followed by an object*

A procedure with functional cohesion usually performs an operation on an object. The name should reflect what the procedure does, and an operation on an object implies a verb-plus-object name. *PrintDocument()*, *CalcMonthlyRevenues()*, *CheckOrderInfo()*, and *RepaginateDocument()* are samples of good procedure names.

In object-oriented languages, you don’t need to include the name of the object in the procedure name because the object itself is included in the call. You invoke routines with statements like *document.Print()*, *orderInfo.Check()*, and *monthlyRevenues.Calc()*. Names like *document.PrintDocument()* are redundant and can become inaccurate when they’re carried through to derived classes. If *Check* is a class derived from *Document*, *check.Print()* seems clearly to be printing a check, whereas *check.PrintDocument()* sounds like it might be printing a checkbook register or monthly statement—but it doesn’t sound like it’s printing a check.

### *Use opposites precisely*

Using naming conventions for opposites helps consistency, which helps readability. Opposite-pairs like *first/last* are commonly understood. Opposite-pairs like *FileOpen()* and *\_lclose()* (from the Windows 3.1 software developer’s kit) are not symmetrical and are confusing. Here are some common opposites:

- add/remove
- begin/end
- create/destroy
- first/last
- get/put
- get/set
- increment/decrement
- insert/delete
- lock/unlock
- min/max
- next/previous

- 462           • old/new
- 463           • open/close
- 464           • show/hide
- 465           • source/target
- 466           • start/stop
- 467           • up/down

### ***Establish conventions for common operations***

In some systems, it's important to distinguish among different kinds of operations. A naming convention is often the easiest and most reliable way of indicating these distinctions.

The code on one of my projects assigned each object a unique identifier. We neglected to establish a convention for naming the routines that would return the object identifier, so we had routine names like these:

```
employee.id.Get()
dependent.GetId()
supervisor()
candidate.id()
```

The *Employee* class exposed its *id* object, which in turn exposed its *Get()* routine. The *Dependent* class exposed a *GetId()* routine. The *Supervisor* class made the *id* its default return value. The *Candidate* class made use of the fact that the *id* object's default return value was the *id*, and exposed the *id* object. By the middle of the project, no one could remember which of these routines was supposed to be used on which object, but by that time too much code had been written to go back and make everything consistent. Consequently, every person on the team had to devote an unnecessary amount of gray matter to remembering the inconsequential detail of which syntax was used on which class to retrieve the *id*. A naming convention for retrieving *ids* would have eliminated this annoyance.

## **7.4 How Long Can a Routine Be?**

On their way to America, the Pilgrims argued about the best maximum length for a routine. After arguing about it for the entire trip, they arrived at Plymouth Rock and started to draft the Mayflower Compact. They still hadn't settled the maximum-length question, and since they couldn't disembark until they'd signed the compact, they gave up and didn't include it. The result has been an interminable debate ever since about how long a routine can be.

497 The theoretical best maximum length is often described as one or two pages of  
498 program listing, 66 to 132 lines. In this spirit, IBM once limited routines to 50  
499 lines, and TRW limited them to two pages (McCabe 1976). Modern programs  
500 tend to have volumes of extremely short routines mixed in with a few longer  
501 routines. Long routines are far from extinct, however. In the Spring of 2003, I  
502 visited two client sites within a month. Programmers at one site were wrestling  
503 with a routine that was about 4,000 lines of code long, and programmers at the  
504 other site were trying to tame a routine that was more than 12,000 lines long!

505 A mountain of research on routine length has accumulated over the years, some  
506 of which is applicable to modern programs, and some of which isn't:

#### 507 **HARD DATA**

- 508 ● A study by Basili and Perricone found that routine size was inversely  
509 correlated with errors; as the size of routines increased (up to 200 lines of  
510 code), the number of errors per line of code decreased (Basili and Perricone  
1984).
- 511 ● Another study found that routine size was not correlated with errors, even  
512 though structural complexity and amount of data were correlated with errors  
513 (Shen et al. 1985).
- 514 ● A 1986 study found that small routines (32 lines of code or fewer) were not  
515 correlated with lower cost or fault rate (Card, Church, and Agresti 1986;  
516 Card and Glass 1990). The evidence suggested that larger routines (65 lines  
517 of code or more) were cheaper to develop per line of code.
- 518 ● An empirical study of 450 routines found that small routines (those with  
519 fewer than 143 source statements, including comments) had 23 percent more  
520 errors per line of code than larger routines but were 2.4 times less expensive  
521 to fix than larger routines (Selby and Basili 1991).
- 522 ● Another study found that code needed to be changed least when routines  
523 averaged 100 to 150 lines of code (Lind and Vairavan 1989).
- 524 ● A study at IBM found that the most error-prone routines were those that  
525 were larger than 500 lines of code. Beyond 500 lines, the error rate tended to  
526 be proportional to the size of the routine (Jones 1986a).

527 Where does all this leave the question of routine length in object-oriented  
528 programs? A large percentage of routines in object-oriented programs will be  
529 accessor routines, which will be very short. From time to time, a complex  
530 algorithm will lead to a longer routine, and in those circumstances, the routine  
531 should be allowed to grow organically up to 100-200 lines. (A line is a  
532 noncomment, nonblank line of source code.) Decades of evidence say that  
533 routines of such length are no more error prone than shorter routines. Let issues  
534 such as depth of nesting, number of variables, and other complexity-related

535

536

537

538

539

540

considerations dictate the length of the routine rather than imposing a length restriction per se.

If you want to write routines longer than about 200 lines, be careful. None of the studies that reported decreased cost, decreased error rates, or both with larger routines distinguished among sizes larger than 200 lines, and you're bound to run into an upper limit of understandability as you pass 200 lines of code.

541

7.5 How to Use Routine Parameters

542

543

544

545

**HARD DATA**

Interfaces between routines are some of the most error-prone areas of a program. One often-cited study by Basili and Perricone (1984) found that 39 percent of all errors were internal interface errors—errors in communication between routines. Here are a few guidelines for minimizing interface problems:

546

*Put parameters in input-modify-output order*

547

548

549

550

551

**CROSS-REFERENCE** For details on documenting routine parameters, see “Commenting Routines” in Section 32.5. For details on formatting parameters, see Section 31.7, “Laying Out Routines.”

Instead of ordering parameters randomly or alphabetically, list the parameters that are input-only first, input-and-output second, and output-only third. This ordering implies the sequence of operations happening within the routine—inputting data, changing it, and sending back a result. Here are examples of parameter lists in Ada:

552

**Ada Example of Parameters in Input-Modify-Output Order**

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

*Ada uses in and out keywords to make input and output parameters clear.*

```
procedure InvertMatrix(  
    originalMatrix: in Matrix;  
    resultMatrix: out Matrix  
);  
...  
  
procedure ChangeSentenceCase(  
    desiredCase: in StringCase;  
    sentence: in out Sentence  
);  
...  
  
procedure PrintPageNumber(  
    pageNumber: in Integer;  
    status: out StatusType  
);
```

569

570

571

572

This ordering convention conflicts with the C-library convention of putting the modified parameter first. The input-modify-output convention makes more sense to me, but if you consistently order parameters in some way, you still do the readers of your code a service.



573 **Create your own *in* and *out* keywords**

574 Other modern languages don't support the *in* and *out* keywords like Ada does. In  
575 those languages, you might still be able to use the preprocessor to create your  
576 own *in* and *out* keywords. Here's how that could be done in C++:

577 **C++ Example of Defining Your Own *In* and *Out* Keywords**

```
578 #define IN
579 #define OUT
580
581 void InvertMatrix(
582     IN Matrix originalMatrix,
583     OUT Matrix *resultMatrix
584 );
585 ...
586
587 void ChangeSentenceCase(
588     IN StringCase desiredCase,
589     IN OUT Sentence *sentenceToEdit
590 );
591 ...
592
593 void PrintPageNumber(
594     IN int pageNumber,
595     OUT StatusType &status
596 );
```

597 In this case, the *IN* and *OUT* macro-keywords are used for documentation  
598 purposes. To make the value of a parameter changeable by the called routine, the  
599 parameter still needs to be passed as a pointer or as a reference parameter.

600 **If several routines use similar parameters, put the similar parameters in a**  
601 **consistent order**

602 The order of routine parameters can be a mnemonic, and inconsistent order can  
603 make parameters hard to remember. For example, in C, the *fprintf()* routine is the  
604 same as the *printf()* routine except that it adds a file as the first argument. A  
605 similar routine, *fputs()*, is the same as *puts()* except that it adds a file as the last  
606 argument. This is an aggravating, pointless difference that makes the parameters  
607 of these routines harder to remember than they need to be.

608 On the other hand, the routine *strncpy()* in C takes the arguments target string,  
609 source string, and maximum number of bytes, in that order, and the routine  
610 *memcpy()* takes the same arguments in the same order. The similarity between  
611 the two routines helps in remembering the parameters in either routine.

612 In Microsoft Windows programming, most of the Windows routines take a  
 613 “handle” as their first parameter. The convention is easy to remember and makes  
 614 each routine’s argument list easier to remember.

615 ***Use all the parameters***

616 **HARD DATA**  
 617 If you pass a parameter to a routine, use it. If you aren’t using it, remove the  
 618 parameter from the routine interface. Unused parameters are correlated with an  
 619 increased error rate. In one study, 46 percent of routines with no unused  
 620 variables had no errors. Only 17 to 29 percent of routines with more than one  
 unreferenced variable had no errors (Card, Church, and Agresti 1986).

621 This rule to remove unused parameters has two exceptions. First, if you’re using  
 622 function pointers in C++, you’ll have several routines with identical parameter  
 623 lists. Some of the routines might not use all the parameters. That’s OK. Second,  
 624 if you’re compiling part of your program conditionally, you might compile out  
 625 parts of a routine that use a certain parameter. Be nervous about this practice, but  
 626 if you’re convinced it works, that’s OK too. In general, if you have a good  
 627 reason not to use a parameter, go ahead and leave it in place. If you don’t have a  
 628 good reason, make the effort to clean up the code.

629 ***Put status or error variables last***

630 By convention, status variables and variables that indicate an error has occurred  
 631 go last in the parameter list. They are incidental to the main purpose of the  
 632 routine, and they are output-only parameters, so it’s a sensible convention.

633 ***Don’t use routine parameters as working variables***

634 It’s dangerous to use the parameters passed to a routine as working variables.  
 635 Use local variables instead. For example, in the Java fragment below, the  
 636 variable *inputVal* is improperly used to store intermediate results of a  
 637 computation.

#### 638 **Java Example of Improper Use of Input Parameters**

```
639 int Sample( int inputVal ) {
640     inputVal = inputVal * CurrentMultiplier( inputVal );
641     inputVal = inputVal + CurrentAdder( inputVal );
642     ...
643     return inputVal;
644 }
```

643 *At this point, inputVal no*  
 644 *longer contains the value that*  
 645 *was input.*

646 *inputVal* in this code fragment is misleading because by the time execution  
 647 reaches the last line, *inputVal* no longer contains the input value; it contains a  
 648 computed value based in part on the input value, and it is therefore misnamed. If  
 649 you later need to modify the routine to use the original input value in some other  
 650 place, you’ll probably use *inputVal* and assume that it contains the original input  
 value when it actually doesn’t.

How do you solve the problem? Can you solve it by renaming *inputVal*? Probably not. You could name it something like *workingVal*, but that's an incomplete solution because the name fails to indicate that the variable's original value comes from outside the routine. You could name it something ridiculous like *InputValThatBecomesWorkingVal* or give up completely and name it *X* or *Val*, but all these approaches are weak.

A better approach is to avoid current and future problems by using working variables explicitly. The following code fragment demonstrates the technique:

#### Java Example of Good Use of Input Parameters

```
int Sample( int inputVal ) {
    int workingVal = inputVal;
    workingVal = workingVal * CurrentMultiplier( workingVal );
    workingVal = workingVal + CurrentAdder( workingVal );
    ...
    ...
    return workingVal;
}
```

If you need to use the original value of *inputVal* here or somewhere else, it's still available.

Introducing the new variable *workingVal* clarifies the role of *inputVal* and eliminates the chance of erroneously using *inputVal* at the wrong time. (Don't take this reasoning as a justification for literally naming a variable *workingVal*. In general, *workingVal* is a terrible name for a variable, and the name is used in this example only to make the variable's role clear.)

Assigning the input value to a working variable emphasizes where the value comes from. It eliminates the possibility that a variable from the parameter list will be modified accidentally. In C++, this practice can be enforced by the compiler using the keyword *const*. If you designate a parameter as *const*, you're not allowed to modify its value within a routine.

**CROSS-REFERENCE** For details on interface assumptions, see the introduction to Chapter 8, "Defensive Programming." For details on documentation, see Chapter 32, "Self-Documenting Code."

#### *Document interface assumptions about parameters*

If you assume the data being passed to your routine has certain characteristics, document the assumptions as you make them. It's not a waste of effort to document your assumptions both in the routine itself and in the place where the routine is called. Don't wait until you've written the routine to go back and write the comments—you won't remember all your assumptions. Even better than commenting your assumptions, use assertions to put them into code.

What kinds of interface assumptions about parameters should you document?

- Whether parameters are input-only, modified, or output-only
- Units of numeric parameters (inches, feet, meters, and so on)

- Meanings of status codes and error values if enumerated types aren't used
- Ranges of expected values
- Specific values that should never appear

---

**HARD DATA*****Limit the number of a routine's parameters to about seven***

Seven is a magic number for people's comprehension. Psychological research has found that people generally cannot keep track of more than about seven chunks of information at once (Miller 1956). This discovery has been applied to an enormous number of disciplines, and it seems safe to conjecture that most people can't keep track of more than about seven routine parameters at once.

In practice, how much you can limit the number of parameters depends on how your language handles complex data types. If you program in a modern language that supports structured data, you can pass a composite data type containing 13 fields and think of it as one mental "chunk" of data. If you program in a more primitive language, you might need to pass all 13 fields individually.

If you find yourself consistently passing more than a few arguments, the coupling among your routines is too tight. Design the routine or group of routines to reduce the coupling. If you are passing the same data to many different routines, group the routines into a class and treat the frequently used data as class data.

***Consider an input, modify, and output naming convention for parameters***

If you find that it's important to distinguish among input, modify, and output parameters, establish a naming convention that identifies them. You could prefix them with *i\_*, *m\_*, and *o\_*. If you're feeling verbose, you could prefix them with *Input\_*, *Modify\_*, and *Output\_*.

***Pass the variables or objects that the routine needs to maintain its interface abstraction***

There are two competing schools of thought about how to pass parameters from an object to a routine. Suppose you have an object that exposes data through 10 access routines, and the called routine needs 3 of those data elements to do its job.

Proponents of the first school of thought argue that only the 3 specific elements needed by the routine should be passed. They argue that that will keep the connections between routines to a minimum, reduce coupling, and make them easier to understand, easier to reuse, and so on. They say that passing the whole object to a routine violates the principle of encapsulation by potentially exposing all 10 access routines to the routine that's called.

Proponents of the second school argue that the whole object should be passed. They argue that the interface can remain more stable if the called routine has the flexibility to use additional members of the object without changing the routine's interface. They argue that passing 3 specific elements violates encapsulation by exposing which specific data elements the routine is using.

I think both these rules are simplistic and miss the most important consideration, which is, *what abstraction is presented by the routine's interface?*

- If the abstraction is that the routine expects you to have 3 specific data elements, and it is only a coincidence that those 3 elements happen to be provided by the same object, then you should pass the 3 specific data elements individually.
- If the abstraction is that you will always have that particular object in hand and the routine will do something or other with that object, then you truly do break the abstraction when you expose the three specific data elements.

If you're passing the whole object and you find yourself creating the object, populating it with the 3 elements needed by the called routine, and then pulling those elements out of the object after the routine is called, that's an indication that you should be passing the 3 specific elements rather than the whole object. (Generally code that "sets up" for a call to a routine or "takes down" after a call to a routine is an indication that the routine is not well designed.)

If you find yourself frequently changing the parameter list to the routine, with the parameters coming from the same object each time, that's an indication that you should be passing the whole object rather than specific elements.

### *Used named parameters*

In some languages, you can explicitly associate formal parameters with actual parameters. This makes parameter usage more self-documenting and helps avoid errors from mismatching parameters. Here's an example in Visual Basic:

---

#### **Visual Basic Example of Explicitly Identifying Parameters**

Here's where the formal parameters are declared.

```
Private Function Distance3d( _
    ByVal xDistance As Coordinate, _
    ByVal yDistance As Coordinate, _
    ByVal zDistance As Coordinate _
)
    ...
End Function
...
Private Function Velocity( _
    ByVal latitude as Coordinate, _
```

```

763         ByVal longitude as Coordinate, _
764         ByVal elevation as Coordinate _
765     )
766     ...
767     Distance = Distance3d( xDistance := latitude, yDistance := longitude, _
768         zDistance := elevation )
769     ...
770 End Function

```

Here's where the actual parameters are mapped to the formal parameters.

This technique is especially useful when you have longer-than-average lists of identically typed arguments, which increases the chances that you can insert a parameter mismatch without the compiler detecting it. Explicitly associating parameters may be overkill in many environments, but in safety-critical or other high-reliability environments the extra assurance that parameters match up the way you expect can be worthwhile.

### ***Don't assume anything about the parameter-passing mechanism***

Some hard-core nanosecond scrapers worry about the overhead associated with passing parameters and bypass the high-level language's parameter-passing mechanism. This is dangerous and makes code nonportable. Parameters are commonly passed on a system stack, but that's hardly the only parameter-passing mechanism that languages use. Even with stack-based mechanisms, the parameters themselves can be passed in different orders and each parameter's bytes can be ordered differently. If you fiddle with parameters directly, you virtually guarantee that your program won't run on a different machine.

### ***Make sure actual parameters match formal parameters***

Formal parameters, also known as dummy parameters, are the variables declared in a routine definition. Actual parameters are the variables or constants used in the actual routine calls.

A common mistake is to put the wrong type of variable in a routine call—for example, using an integer when a floating point is needed. (This is a problem only in weakly typed languages like C when you're not using full compiler warnings. Strongly typed languages such as C++ and Java don't have this problem.) When arguments are input only, this is seldom a problem; usually the compiler converts the actual type to the formal type before passing it to the routine. If it is a problem, usually your compiler gives you a warning. But in some cases, particularly when the argument is used for both input and output, you can get stung by passing the wrong type of argument.

Develop the habit of checking types of arguments in parameter lists and heeding compiler warnings about mismatched parameter types.

## 7.6 Special Considerations in the Use of Functions

Modern languages such as C++, Java, and Visual Basic support both functions and procedures. A function is a routine that returns a value; a procedure is a routine that does not. This distinction is as much a semantic distinction as a syntactic one. In C++, all routines are typically called “functions,” however, a function with a *void* return type is semantically a procedure and should be treated as such.

### When to Use a Function and When to Use a Procedure

Purists argue that a function should return only one value, just as a mathematical function does. This means that a function would take only input parameters and return its only value through the function itself. The function would always be named for the value it returned, as *sin()*, *CustomerID()*, and *ScreenHeight()* are. A procedure, on the other hand, could take input, modify, and output parameters—as many of each as it wanted to.

A common programming practice is to have a function that operates as a procedure and returns a status value. Logically, it works as a procedure, but because it returns a value, it’s officially a function. For example, you might have a routine called *FormatOutput()* used with a *report* object in statements like this one:

```
if ( report.FormatOutput( formattedReport ) = Success ) then ...
```

In this example, *report.FormatOutput()* operates as a procedure in that it has an output parameter, *formattedReport*, but it is technically a function because the routine itself returns a value. Is this a valid way to use a function? In defense of this approach, you could maintain that the function return value has nothing to do with the main purpose of the routine, formatting output, or with the routine name, *report.FormatOutput()*; in that sense it operates more as a procedure does even if it is technically a function. The use of the return value to indicate the success or failure of the procedure is not confusing if the technique is used consistently.

The alternative is to create a procedure that has a status variable as an explicit parameter, which promotes code like this fragment:

```
report.FormatOutput( formattedReport, outputStatus )  
if ( outputStatus = Success ) then ...
```

836 I prefer the second style of coding, not because I'm hard-nosed about the  
 837 difference between functions and procedures but because it makes a clear  
 838 separation between the routine call and the test of the status value. To combine  
 839 the call and the test into one line of code increases the density of the statement  
 840 and correspondingly its complexity. The following use of a function is fine too:

```
841         outputStatus = report.FormatOutput( formattedReport )
842         if ( outputStatus = Success ) then ...
```

#### 843 **KEY POINT**

844 In short, use a function if the primary purpose of the routine is to return the value indicated by the function name. Otherwise, use a procedure.

## 845 **Setting the Function's Return Value**

846 Using a function creates the risk that the function will return an incorrect return  
 847 value. This usually happens when the function has several possible paths and one  
 848 of the paths doesn't set a return value.

### 849 *Check all possible return paths*

850 When creating a function, mentally execute each path to be sure that the function  
 851 returns a value under all possible circumstances. It's good practice to initialize  
 852 the return value at the beginning of the function to a default value—which  
 853 provides a safety net in the event of that the correct return value is not set.

### 854 *Don't return references or pointers to local data*

855 As soon as the routine ends and the local data goes out of scope, the reference or  
 856 pointer to the local data will be invalid. If an object needs to return information  
 857 about its internal data, it should save the information as class member data. It  
 858 should then provide accessor functions that return the values of the member data  
 859 items rather than references or pointers to local data.

## 860 **7.7 Macro Routines and Inline Routines**

861 **CROSS-REFERENCE** Even  
 862 n if your language doesn't  
 863 have a macro preprocessor,  
 864 you can build your own. For  
 details, see Section 30.5,  
 "Building Your Own  
 865 Programming Tools."

866 Routines created with preprocessor macros call for a few unique considerations.  
 867 The following rules and examples pertain to using the preprocessor in C++. If  
 868 you're using a different language or preprocessor, adapt the rules to your  
 situation.

### 865 *Fully parenthesize macro expressions*

866 Because macros and their arguments are expanded into code, be careful that they  
 867 expand the way you want them to. One common problem lies in creating a  
 868 macro like this one:



---

**C++ Example of a Macro That Doesn't Expand Properly**

```
#define Cube( a ) a*a*a
```

This macro has a problem. If you pass it nonatomic values for  $a$ , it won't do the multiplication properly. If you use the expression `Cube(  $x+1$  )`, it expands to  $x+1 * x + 1 * x + 1$ , which, because of the precedence of the multiplication and addition operators, is not what you want. A better but still not perfect version of the macro looks like this:

---

**C++ Example of a Macro That Still Doesn't Expand Properly**

```
#define Cube( a ) (a)*(a)*(a)
```

This is close, but still no cigar. If you use `Cube()` in an expression that has operators with higher precedence than multiplication, the  $(a)*(a)*(a)$  will be torn apart. To prevent that, enclose the whole expression in parentheses:

---

**C++ Example of a Macro That Works**

```
#define Cube( a ) ((a)*(a)*(a))
```

---

***Surround multiple-statement macros with curly braces***

A macro can have multiple statements, which is a problem if you treat it as if it were a single statement. Here's an example of a macro that's headed for trouble:

---

**C++ Example of a Macro with Multiple Statements That Doesn't Work**

```
#define LookupEntry( key, index ) \
    index = (key - 10) / 5; \
    index = min( index, MAX_INDEX ); \
    index = max( index, MIN_INDEX ); \
    ...

for ( entryCount = 0; entryCount < numEntries; entryCount++ )
    LookupEntry( entryCount, tableIndex[ entryCount ] );
```

This macro is headed for trouble because it doesn't work as a regular function would. As it's shown, the only part of the macro that's executed in the *for* loop is the first line of the macro:

```
    index = (key - 10) / 5;
```

To avoid this problem, surround the macro with curly braces, as shown here:

---

**C++ Example of a Macro with Multiple Statements That Works**

```
#define LookupEntry( key, index ) { \
    index = (key - 10) / 5; \
    index = min( index, MAX_INDEX ); \
    index = max( index, MIN_INDEX ); \
}
```

The practice of using macros as substitutes for function calls is generally considered risky and hard to understand—bad programming practice—so use this technique only if your specific circumstances require it.

***Name macros that expand to code like routines so that they can be replaced by routines if necessary***

The C++-language convention for naming macros is to use all capital letters. If the macro can be replaced by a routine, however, name it using the naming convention for routines instead. That way you can replace macros with routines and vice versa without changing anything but the routine involved.

Following this recommendation entails some risk. If you commonly use ++ and -- as side effects (as part of other statements), you'll get burned when you use macros that you think are routines. Considering the other problems with side effects, this is just one more reason to avoid using side effects.

## Limitations on the Use of Macro Routines

Modern languages like C++ provide numerous alternatives to the use of macros:

- *const* for declaring constant values
- *inline* for defining functions that will be compiled as inline code
- *template* for defining standard operations like *min*, *max*, and so on in a type-safe way
- *enum* for defining enumerated types
- *typedef* for defining simple type substitutions

### KEY POINT

As Bjarne Stroustrup, designer of C++ points out, “Almost every macro demonstrates a flaw in the programming language, in the program, or in the programmer.... When you use macros, you should expect inferior service from tools such as debuggers, cross-reference tools, and profilers” (Stroustrup 1997). Macros are useful for supporting conditional compilation (see Section 8.6), but careful programmers generally use a macro as an alternative to a routine only as a last resort.

## Inline Routines

C++ supports an *inline* keyword. An *inline* routine allows the programmer to treat the code as a routine at code-writing time. But the compiler will convert each instance of the routine into inline code at compile time. The theory is that *inline* can help produce highly efficient code that avoids routine-call overhead.

939

940

941

942

943

944

945

946

947

948

949

950

*Use inline routines sparingly*

Inline routines violate encapsulation because C++ requires the programmer to put the code for the implementation of the inline routine in the header file, which exposes it to every programmer who uses the header file.

Inline routines require a routine’s full code to be generated every time the routine is invoked, which for an inline routine of any size will increase code size. That can create problems of its own.

The bottom line on inlining for performance reasons is the same as the bottom line on any other coding technique that’s motivated by performance—profile the code and measure the improvement. If the anticipated performance gain doesn’t justify the bother of profiling the code to verify the improvement, it doesn’t justify the erosion in code quality either.

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

**CROSS-REFERENCE** This is a checklist of considerations about the quality of the routine. For a list of the steps used to build a routine, see the checklist “The Pseudocode Programming Process” in Chapter 9, page 000.

---

**CHECKLIST: High-Quality Routines**

---

**Big-Picture Issues**

☐

Is the reason for creating the routine sufficient?

☐

Have all parts of the routine that would benefit from being put into routines of their own been put into routines of their own?

☐

Is the routine’s name a strong, clear verb-plus-object name for a procedure or a description of the return value for a function?

☐

Does the routine’s name describe everything the routine does?

☐

Have you established naming conventions for common operations?

☐

Does the routine have strong, functional cohesion—doing one and only one thing and doing it well?

☐

Do the routines have loose coupling—are the routine’s connections to other routines small, intimate, visible, and flexible?

☐

Is the length of the routine determined naturally by its function and logic, rather than by an artificial coding standard?

**Parameter-Passing Issues**

☐

Does the routine’s parameter list, taken as a whole, present a consistent interface abstraction?

☐

Are the routine’s parameters in a sensible order, including matching the order of parameters in similar routines?

☐

Are interface assumptions documented?

☐

Does the routine have seven or fewer parameters?

☐

Is each input parameter used?

- 974            ☐ Is each output parameter used?
- 975            ☐ Does the routine avoid using input parameters as working variables?
- 976            ☐ If the routine is a function, does it return a valid value under all possible
- 977                                  circumstances?
- 978
- 

979

## Key Points

- 980            • The most important reason to create a routine is to improve the intellectual
- 981                                  manageability of a program, and you can create a routine for many other
- 982                                  good reasons. Saving space is a minor reason; improved readability,
- 983                                  reliability, and modifiability are better reasons.
- 984            • Sometimes the operation that most benefits from being put into a routine of
- 985                                  its own is a simple one.
- 986            • The name of a routine is an indication of its quality. If the name is bad and
- 987                                  it's accurate, the routine might be poorly designed. If the name is bad and
- 988                                  it's inaccurate, it's not telling you what the program does. Either way, a bad
- 989                                  name means that the program needs to be changed.
- 990            • Functions should be used only when the primary purpose of the function is
- 991                                  to return the specific value described by the function's name.
- 992            • Careful programmers use macro routines and inline routines with care, and
- 993                                  only as a last resort.