

12

Fundamental Data Types

Contents

- 12.1 Numbers in General
- 12.2 Integers
- 12.3 Floating-Point Numbers
- 12.4 Characters and Strings
- 12.5 Boolean Variables
- 12.6 Enumerated Types
- 12.7 Named Constants
- 12.8 Arrays
- 12.9 Creating Your Own Types

Related Topics

- Naming data: Chapter 11
- Unusual data types: Chapter 13
- General issues in using variables: Chapter 10
- Formatting data declarations: “Laying Out Data Declarations” in Section 31.5
- Documenting variables: “Commenting Data Declarations” in Section 32.5
- Creating classes: Chapter 6

THE FUNDAMENTAL DATA TYPES ARE the basic building blocks for all other data types. This chapter contains tips for using integers, floating-point numbers, characters and strings, boolean variables, enumerated types, named constants, and arrays. The final section in this chapter describes how to create your own types.

This chapter covers basic troubleshooting for the fundamental types of data. If you’ve got your fundamental-data bases covered, skip to the end of the chapter, review the checklist of problems to avoid, and move on to the discussion of unusual data types in Chapter 13.

12.1 Numbers in General

Here are several guidelines for making your use of numbers less error prone.

Avoid “magic numbers.”

Magic numbers are literal numbers such as *100* or *47524* that appear in the middle of a program without explanation. If you program in a language that supports named constants, use them instead. If you can’t use named constants, use global variables when it is feasible to.

Avoiding magic numbers yields three advantages:

- Changes can be made more reliably. If you use named constants, you won’t overlook one of the *100*s, or change a *100* that refers to something else.
- Changes can be made more easily. When the maximum number of entries changes from *100* to *200*, if you’re using magic numbers you have to find all the *100*s and change them to *200*s. If you use *100+1* or *100-1* you’ll also have to find all the *101*s and *99*s and change them to *201*s and *199*s. If you’re using a named constant, you simply change the definition of the constant from *100* to *200* in one place.
- Your code is more readable. Sure, in the expression
for i = 0 to 99 do ...
you can guess that *99* refers to the maximum number of entries. But the expression
for i = 0 to MAX_ENTRIES-1 do ...
leaves no doubt. Even if you’re certain that a number will never change, you get a readability benefit if you use a named constant.

Use hard-coded 0s and 1s if you need to

The values *0* and *1* are used to increment, decrement, and start loops at the first element of an array. The *0* in

for i = 0 to CONSTANT do ...
is OK, and the *1* in

total = total + 1
is OK. A good rule of thumb is that the only literals that should occur in the body of a program are *0* and *1*. Any other literals should be replaced with something more descriptive.

61 *Anticipate divide-by-zero errors*

62 Each time you use the division symbol (/ in most languages), think about
 63 whether it's possible for the denominator of the expression to be 0. If the
 64 possibility exists, write code to prevent a divide-by-zero error.

65 *Make type conversions obvious*

66 Make sure that someone reading your code will be aware of it when a conversion
 67 between different data types occurs. In C++ you could say

68 `y = x + (float) i`

69 and in Visual Basic you could say

70 `y = x + CSng(i)`

71 This practice also helps to ensure that the conversion is the one you want to
 72 occur—different compilers do different conversions, so you're taking your
 73 chances otherwise.

74 **CROSS-REFERENCE** For
 75 a variation on this example,
 76 see "Avoid equality
 77 comparisons" in Section
 78 12.3.

Avoid mixed-type comparisons

If x is a floating-point number and i is an integer, the test

`if (i = x) ...`

is almost guaranteed not to work. By the time the compiler figures out which
 type it wants to use for the comparison, converts one of the types to the other,
 does a bunch of rounding, and determines the answer, you'll be lucky if your
 program runs at all. Do the conversion manually so that the compiler can
 compare two numbers of the same type and you know exactly what's being
 compared.

83 **KEY POINT**

Heed your compiler's warnings

84 Many modern compilers tell you when you have different numeric types in the
 85 same expression. Pay attention! Every programmer has been asked at one time or
 86 another to help someone track down a pesky error, only to find that the compiler
 87 had warned about the error all along. Top programmers fix their code to
 88 eliminate all compiler warnings. It's easier to let the compiler do the work than
 89 to do it yourself.

90 **12.2 Integers**

91 Here are a few considerations to bear in mind when using integers:

Check for integer division

92 When you're using integers, $7/10$ does not equal 0.7. It usually equals 0. This
 93 applies equally to intermediate results. In the real world $10 * (7/10) = (10*7) / 10$
 94 $= 7$. Not so in the world of integer arithmetic. $10 * (7/10)$ equals 0 because the
 95

integer division (7/10) equals 0. The easiest way to remedy this problem is to reorder the expression so that the divisions are done last: $(10*7) / 10$.

Check for integer overflow

When doing integer multiplication or addition, you need to be aware of the largest possible integer. The largest possible unsigned integer is often 65,535, or $2^{32}-1$. The problem comes up when you multiply two numbers that produce a number bigger than the maximum integer. For example, if you multiply 250 * 300, the right answer is 75,000. But if the maximum integer is 65,535, the answer you'll get is probably 9464 because of integer overflow (75,000 - 65,536 = 9464). Here are the ranges of common integer types:

Integer Type	Range
Signed 8-bit	-128 through 127
Unsigned 8-bit	0 through 255
Signed 16-bit	-32,768 through 32,767
Unsigned 16-bit	0 through 65,535
Signed 32-bit	-2,147,483,648 through 2,147,483,647
Unsigned 32-bit	0 through 4,294,967,295
Signed 64-bit	-9,223,372,036,854,775,808 through 9,223,372,036,854,775,807
Unsigned 64-bit	0 through 18,446,744,073,709,551,615

The easiest way to prevent integer overflow is to think through each of the terms in your arithmetic expression and try to imagine the largest value each can assume. For example, if in the integer expression $m = j * k$, the largest expected value for j is 200 and the largest expected value for k is 25, the largest value you can expect for m is $200 * 25 = 5,000$. This is OK on a 32-bit machine since the largest integer is 2,147,483,647. On the other hand, if the largest expected value for j is 200,000 and the largest expected value for k is 100,000, the largest value you can expect for m is $200,000 * 100,000 = 20,000,000,000$. This is not OK since 20,000,000,000 is larger than 2,147,483,647. In this case, you would have to use 64-bit integers or floating-point numbers to accommodate the largest expected value of m .

Also consider future extensions to the program. If m will never be bigger than 5,000, that's great. But if you expect m to grow steadily for several years, take that into account.

Check for overflow in intermediate results

The number at the end of the equation isn't the only number you have to worry about. Suppose you have the following code:

123
124
125
126
127
128
129
130
131
132

Java Example of Overflow of Intermediate Results

```
int termA = 1000000;
int termB = 1000000;
int product = termA * termB / 1000000;
System.out.println( "( " + termA + " * " + termB + " ) / 1000000 = " + product );
```

If you think the *Product* assignment is the same as $(100,000 * 100,000) / 100,000$, you might expect to get the answer *100,000*. But the code has to compute the intermediate result of $100,000 * 100,000$ before it can divide by the final *100,000*, and that means it needs a number as big as *1,000,000,000,000*. Guess what? Here's the result:

133
134
135
136
137

(1000000 * 1000000) / 1000000 = -727

If your integers go to only 2,147,483,647, the intermediate result is too large for the integer data type. In this case, the intermediate result that should be *1,000,000,000,000* is 727,379,968, so when you divide by *100,000*, you get -727, rather than *100,000*.

138
139

You can handle overflow in intermediate results the same way you handle integer overflow, by switching to a long-integer or floating-point type.

140

12.3 Floating-Point Numbers

KEY POINT

141
142
143
144
145
146
147

The main consideration in using floating-point numbers is that many fractional decimal numbers can't be represented accurately using the 1s and 0s available on a digital computer. Nonterminating decimals like 1/3 or 1/7 can usually be represented to only 7 or 15 digits of accuracy. In my version of Visual Basic, a 32 bit floating-point representation of 1/3 equals 0.33333330. It's accurate to 7 digits. This is accurate enough for most purposes, but inaccurate enough to trick you sometimes.

148

Here are a few specific guidelines for using floating-point numbers:

149
150
151
152
153
154

Avoid additions and subtractions on numbers that have greatly different magnitudes

With a 32-bit floating-point variable, $1,000,000.00 + 0.1$ probably produces an answer of 1,000,000.00 because 32 bits don't give you enough significant digits to encompass the range between 1,000,000 and 0.1. Likewise, $5,000,000.02 - 5,000,000.01$ is probably 0.0.

155 **CROSS-REFERENCE** For
 156 algorithms books that
 157 describe ways to solve these
 158 problems, see “Additional
 159 Resources on Data Types” in
 160 Section 10.1.

161 *1 is equal to 2 for*
 162 *sufficiently large values*
 163 *of 1.*

164 —Anonymous

167
 168 *The variable nominal is a 64-*
 169 *bit real.*

170
 171
 172 *sum is computed as 10*0.1. It*
 173 *should be 1.0.*

174
 175
 176 *Here's the bad comparison.*

177
 178
 179
 180

181 As you can probably guess, the output from this program is

182
 183 Numbers are different.

The line-by-line values of *sum* in the *for* loop look like this:

184 0.1
 185 0.2
 186 0.30000000000000004
 187 0.4
 188 0.5
 189 0.6
 190 0.7
 191 0.7999999999999999
 192 0.8999999999999999
 193 0.9999999999999999

194 Thus, it's a good idea to find an alternative to using an equality comparison for
 195 floating point numbers. One effective approach is to determine a range of
 196 accuracy that is acceptable and then use a boolean function to determine whether

Solutions? If you have to add a sequence of numbers that contains huge differences like this, sort the numbers first, and then add them starting with the smallest values. Likewise, if you need to sum an infinite series, start with the smallest term—essentially, sum the terms backwards. This doesn't eliminate round-off problems, but it minimizes them. Many algorithms books have suggestions for dealing with cases like this.

Avoid equality comparisons

Floating-point numbers that should be equal are not always equal. The main problem is that two different paths to the same number don't always lead to the same number. For example, 0.1 added 10 times rarely equals 1.0. The first example on the next page shows two variables, *nominal* and *sum*, that should be equal but aren't.

Java Example of a Bad Comparison of Floating-Point Numbers

```
double nominal = 1.0;
double sum = 0.0;

for ( int i = 0; i < 10; i++ ) {
    sum += 0.1;
}

if ( nominal == sum ) {
    System.out.println( "Numbers are the same." );
}
else {
    System.out.println( "Numbers are different." );
}
```

As you can probably guess, the output from this program is

Numbers are different.

The line-by-line values of *sum* in the *for* loop look like this:

```
0.1
0.2
0.30000000000000004
0.4
0.5
0.6
0.7
0.7999999999999999
0.8999999999999999
0.9999999999999999
```

Thus, it's a good idea to find an alternative to using an equality comparison for floating point numbers. One effective approach is to determine a range of accuracy that is acceptable and then use a boolean function to determine whether

197 the values are close enough. Typically, you would write an *Equals()* function
198 that returns *True* if the values are close enough and *False* otherwise. In Java,
199 such a function would look like this:

200 **CROSS-REFERENCE** This
example is proof of the
201 maxim that there's an
202 exception to every rule.
203 Variables in this realistic
204 example have digits in their
205 names. For the rule *against*
206 using digits in variable
207 names, see Section 11.7,
208 "Kinds of Names to Avoid."

Java Example of a Routine to Compare Floating-Point Numbers

```
double const ACCEPTABLE_DELTA = 0.00001;  
boolean Equals( double Term1, double Term2 ) {  
    if ( Math.abs( Term1 - Term2 ) < ACCEPTABLE_DELTA ) {  
        return true;  
    }  
    else {  
        return false;  
    }  
}
```

If the code in the "bad comparison of floating-point numbers" example were converted so that this routine could be used for comparisons, the new comparison would look like this:

```
if ( Equals( Nominal, Sum ) ) ...
```

The output from the program when it uses this test is

```
Numbers are the same.
```

Depending on the demands of your application, it might be inappropriate to use a hard-coded value for *AcceptableDelta*. You might need to compute *AcceptableDelta* based on the size of the two numbers being compared.

Anticipate rounding errors

Rounding-error problems are no different from the problem of numbers with greatly different magnitudes. The same issue is involved, and many of the same techniques help to solve rounding problems. In addition, here are common specific solutions to rounding problems:

First, change to a variable type that has greater precision. If you're using single-precision floating point, change to double-precision floating point, and so on.

Second, change to binary coded decimal (BCD) variables. The BCD scheme is typically slower and takes up more storage space but prevents many rounding errors. This is particularly valuable if the variables you're using represent dollars and cents or other quantities that must balance precisely.

Third, change from floating-point to integer variables. This is a roll-your-own approach to BCD variables. You will probably have to use 64-bit integers to get the precision you want. This technique requires you to keep track of the fractional part of your numbers yourself. Suppose you were originally keeping track of dollars using floating point with cents expressed as fractional parts of

dollars. This is a normal way to handle dollars and cents. When you switch to integers, you have to keep track of cents using integers and of dollars using multiples of 100 cents. In other words, you multiply dollars by 100 and keep the cents in the 0-to-99 range of the variable. This might seem absurd at first glance, but it's an effective solution in terms of both speed and accuracy. You can make these manipulations easier by creating a *DollarsAndCents* class that hides the integer representation and supports the necessary numeric operations.

Check language and library support for specific data types

Some languages including Visual Basic have data types such as *Currency* that specifically support data that is sensitive to rounding errors. If your language has a built-in data type that provides such functionality, use it!

12.4 Characters and Strings

Here are some tips for using strings. The first applies to strings in all languages.

Avoid magic characters and strings

Magic characters are literal characters (such as `<,$QS>A<,$QS>`) and magic strings are literal strings (such as `<,$QD>Gigamatic Accounting Program<,$QD>`) that appear throughout a program. If you program in a language that supports the use of named constants, use them instead. Otherwise, use global variables. Several reasons for avoiding literal strings follow.

- For commonly occurring strings like the name of your program, command names, report titles, and so on, you might at some point need to change the string's contents. For example, "*Gigamatic Accounting Program*" might change to "*New and Improved! Gigamatic Accounting Program*" for a later version.
- International markets are becoming increasingly important, and it's easier to translate strings that are grouped in a string resource file than it is to translate to them *in situ* throughout a program.
- String literals tend to take up a lot of space. They're used for menus, messages, help screens, entry forms, and so on. If you have too many, they grow beyond control and cause memory problems. String space isn't a concern in many environments, but in embedded systems programming and other applications in which storage space is at a premium, solutions to string-space problems are easier to implement if the strings are relatively independent of the source code.
- Character and string literals are cryptic. Comments or named constants clarify your intentions. In the example below, the meaning of

271 <;\$QS>\027<;\$QS> isn't clear. The use of the *ESCAPE* constant makes the
 272 meaning more obvious.

273 C++ Examples of Comparisons Using Strings

274 *Bad!* if (input_char == '\027') ...
 275 *Better!* if (input_char == ESCAPE) ...

276 *Watch for off-by-one errors*

277 Because substrings can be indexed much as arrays are, watch for off-by-one
 278 errors that read or write past the end of a string.

279 CC2E.COM/1285 *Know how your language and environment support Unicode*

280 In some languages such as Java, all strings are Unicode. In others such as C and
 281 C++, handling Unicode strings requires its own set of functions. Conversion
 282 between Unicode and other character sets is often required for communication
 283 with standard and third-party libraries. If some strings won't be in Unicode (for
 284 example, in C or C++), decide early on whether to use the Unicode character set
 285 at all. If you decide to use Unicode strings, decide where and when to use them.

286 *Decide on an internationalization/localization strategy early in the lifetime* 287 *of a program*

288 Issues related to internationalization and localization are major issues. Key
 289 considerations are deciding whether to store all strings in an external resource
 290 and whether to create separate builds for each language or to determine the
 291 specific language at run-time.

292 CC2E.COM/1292 *If you know you only need to support a single alphabetic language,* 293 *consider using an ISO 8859 character set*

294 For applications that need to support only a single alphabetic language such as
 295 English, and that don't need to support multiple languages or an ideographic
 296 language such as written Chinese, the ISO 8859 extended-ASCII-type standard
 297 makes a good alternative to Unicode.

298 *If you need to support multiple languages, use Unicode*

299 Unicode provides more comprehensive support for international character sets
 300 than ISO 8859 or other standards.

301 *Decide on a consistent conversion strategy among string types*

302 If you use multiple string types, one common approach that helps keep the string
 303 types distinct is to keep all strings in a single format within the program, and
 304 convert the strings to other formats as close as possible to input and output
 305 operations.

Strings in C

C++'s standard template library string class has eliminated most of the traditional problems with strings in C. For those programmers working directly with C strings, here are some ways to avoid common pitfalls.

Be aware of the difference between string pointers and character arrays

The problem with string pointers and character arrays arises because of the way C handles strings. Be alert to the difference between them in two ways:

- Be suspicious of any expression containing a string that involves an equal sign. String operations in C are nearly always done with *strcmp()*, *strcpy()*, *strlen()*, and related routines. Equal signs often imply some kind of pointer error. In C, assignments do not copy string literals to a string variable. Suppose you have a statement like

```
StringPtr = "Some Text String";
```

In this case, `<$QD>Some Text String<$QD>` is a pointer to a literal text string and the assignment merely sets the pointer *StringPtr* to point to the text string. The assignment does not copy the contents to *StringPtr*.
- Use a naming convention to indicate whether the variables are arrays of characters or pointers to strings. One common convention is to use *ps* as a prefix to indicate a *pointer* to a *string* and *ach* as a prefix for an *array* of *characters*. Although they're not always wrong, you should regard expressions involving both *ps* and *ach* prefixes with suspicion.

Declare C-style strings to have length CONSTANT+1

In C and C++, off-by-one errors with C-style strings are easy to make because it's easy to forget that a string of length *n* requires *n + 1* bytes of storage and to forget to leave room for the null terminator (the byte set to 0 at the end of the string). An easy and effective way to avoid such problems is to use named constants to declare all strings. A key in this approach is that you use the named constant the same way every time. Declare the string to be length *CONSTANT+1*, and then use *CONSTANT* to refer to the length of a string in the rest of the code. Here's an example:

C Example of Good String Declarations

```
/* Declare the string to have length of "constant+1".
   Every other place in the program, "constant" rather
   than "constant+1" is used. */
char string[ NAME_LENGTH + 1 ] = { 0 }; /* string of length NAME_LENGTH */
...
/* Example 1: Set the string to all 'A's using the constant,
   NAME_LENGTH, as the number of 'A's that can be copied.
```

345
 346 *Operations on the string*
 347 NAME_LENGTH here...

348
 349
 350
 351
 352 ...and here.

```
Note that NAME_LENGTH rather than NAME_LENGTH + 1 is used. */
for ( i = 0; i < NAME_LENGTH; i++ )
    string[ i ] = 'A';
...

/* Example 2: Copy another string into the first string using
   the constant as the maximum length that can be copied. */
strncpy( string, some_other_string, NAME_LENGTH );
```

If you don't have a convention to handle this, you'll sometimes declare the string to be of length *NAME_LENGTH* and have operations on it with *NAME_LENGTH-1*; at other times you'll declare the string to be of length *NAME_LENGTH+1* and have operations on it work with length *NAME_LENGTH*. Every time you use a string, you'll have to remember which way you declared it.

359 When you use strings the same way every time, you don't have to remember
 360 how you dealt with each string individually and you eliminate mistakes caused
 361 by forgetting the specifics of an individual string. Having a convention
 362 minimizes mental overload and programming errors.

363 **CROSS-REFERENCE** For
 364 more details on initializing
 365 data, see Section 10.3,
 366 "Guidelines for Initializing
 367 Variables."

Initialize strings to null to avoid endless strings

C determines the end of a string by finding a null terminator, a byte set to 0 at the end of the string. No matter how long you think the string is, C doesn't find the end of the string until it finds a 0 byte. If you forget to put a null at the end of the string, your string operations might not act the way you expect them to.

You can avoid endless strings in two ways. First, initialize arrays of characters to 0 when you declare them, as shown below:

C Example of a Good Declaration of a Character Array

```
char EventName[ MAX_NAME_LENGTH + 1 ] = { 0 };
```

Second, when you allocate strings dynamically, initialize them to 0 by using *calloc()* instead of *malloc()*. *calloc()* allocates memory and initializes it to 0. *malloc()* allocates memory without initializing it so you get potluck when you use memory allocated by *malloc()*.

376 **CROSS-REFERENCE** For
 377 more discussion of arrays,
 378 read Section 12.8, "Arrays,"
 379 later in this chapter.

Use arrays of characters instead of pointers in C

If memory isn't a constraint—and often it is not—declare all your string variables as arrays of characters. This helps to avoid pointer problems, and the compiler will give you more warnings when you do something wrong.

Use strncpy() instead of strcpy() to avoid endless strings

String routines in C come in safe versions and dangerous versions. The more dangerous routines such as *strcpy()* and *strcmp()* keep going until they run into a NULL terminator. Their safer companions, *strncpy()* and *strncmp()*, take a

parameter for maximum length, so that even if the strings go on forever, your function calls won't.

12.5 Boolean Variables

It's hard to misuse logical or boolean variables, and using them thoughtfully makes your program cleaner.

CROSS-REFERENCE For details on using comments to document your program, see Chapter 32, "Self-Documenting Code."

CROSS-REFERENCE For an example of using a boolean function to document your program, see "Making Complicated Expressions Simple" in Section 19.1.

Use boolean variables to document your program

Instead of merely testing a boolean expression, you can assign the expression to a variable that makes the implication of the test unmistakable. For example, in the fragment below, it's not clear whether the purpose of the *if* test is to check for completion, for an error condition, or for something else:

Java Example of Boolean Test in Which the Purpose Is Unclear

```
if ( ( elementIndex < 0 ) || ( MAX_ELEMENTS < elementIndex ) ||  
    ( elementIndex == lastElementIndex )  
    ) {  
    ...  
}
```

In the next fragment, the use of boolean variables makes the purpose of the *if* test clearer:

Java Example of Boolean Test in Which the Purpose Is Clear

```
finished = ( ( elementIndex < 0 ) || ( MAX_ELEMENTS < elementIndex ) );  
repeatedEntry = ( elementIndex == lastElementIndex );  
if ( finished || repeatedEntry ) {  
    ...  
}
```

Use boolean variables to simplify complicated tests

Often when you have to code a complicated test, it takes several tries to get it right. When you later try to modify the test, it can be hard to understand what the test was doing in the first place. Logical variables can simplify the test. In the example above, the program is really testing for two conditions: whether the routine is finished and whether it's working on a repeated entry. By creating the boolean variables *finished* and *repeatedEntry*, you make the *if* test simpler—easier to read, less error prone, and easier to modify.

Here's another example of a complicated test:

Visual Basic Example of a Complicated Test

```
If ( ( document.AtEndOfStream() ) And ( Not inputError ) ) And _
```

CODING HORROR

```

419      ( ( MIN_LINES <= lineCount ) And ( lineCount <= MAX_LINES ) ) And _
420      ( Not ErrorProcessing() ) Then
421      ' do something or other
422      ...
423  End If

```

The test in the example is fairly complicated but not uncommonly so. It places a heavy mental burden on the reader. My guess is that you won't even try to understand the *if* test but will look at it and say, "I'll figure it out later if I really need to." Pay attention to that thought because that's exactly the same thing other people do when they read your code and it contains tests like this.

Here's a rewrite of the code with boolean variables added to simplify the test:

Visual Basic Example of a Simplified Test

Here's the simple test.

```

430  allDataRead = ( document.AtEndOfStream() ) And ( Not inputError )
431  legalLineCount = ( MIN_LINES <= lineCount ) And ( lineCount <= MAX_LINES )
432  If ( allDataRead ) And ( legalLineCount ) And ( Not ErrorProcessing() ) Then
433      ' do something or other
434      ...
435  End If

```

This second version is simpler. My guess is that you'll read the boolean expression in the *if* test without any difficulty.

Create your own boolean type, if necessary

Some languages, such as C++, Java, and Visual Basic have a predefined boolean type. Others, such as C, do not. In languages such as C, you can define your own boolean type. In C, you'd do it this way:

C Example of Defining the BOOLEAN Type

```

443  typedef int BOOLEAN;    // define the boolean type

```

Declaring variables to be *BOOLEAN* rather than *int* makes their intended use more obvious and makes your program a little more self-documenting.

12.6 Enumerated Types

An enumerated type is a type of data that allows each member of a class of objects to be described in English. Enumerated types are available in C++, and Visual Basic and are generally used when you know all the possible values of a variable and want to express them in words. Here are several examples of enumerated types in Visual Basic:

Visual Basic Examples of Enumerated Types

```

453  Public Enum Color

```

```
455     Color_Red
456     Color_Green
457     Color_Blue
458 End Enum
459
460 Public Enum Country
461     Country_China
462     Country_England
463     Country_France
464     Country_Germany
465     Country_India
466     Country_Japan
467     Country_Usa
468 End Enum
469
470 Public Enum Output
471     Output_Screen
472     Output_Printer
473     Output_File
474 End Enum
```

Enumerated types are a powerful alternative to shopworn schemes in which you explicitly say, “1 stands for red, 2 stands for green, 3 stands for blue,...” This ability suggests several guidelines for using enumerated types.

Use enumerated types for readability

Instead of writing statements like

```
480     if chosenColor = 1
481 you can write more readable expressions like
```

```
482     if chosenColor = Color_Red
483 Anytime you see a numeric literal, ask whether it makes sense to replace it with
484 an enumerated type.
```

Use enumerated types for reliability

With a few languages (Ada in particular), an enumerated type lets the compiler perform more thorough type checking than it can with integer values and constants. With named constants, the compiler has no way of knowing that the only legal values are *Color_Red*, *Color_Green*, and *Color_Blue*. The compiler won't object to statements like *color = Country_England* or *country = Output_Printer*. If you use an enumerated type, declaring a variable as *Color*, the compiler will allow the variable to be assigned only the values *Color_Red*, *Color_Green*, and *Color_Blue*.

494 *Use enumerated types for modifiability*

495 Enumerated types make your code easy to modify. If you discover a flaw in your
 496 “1 stands for red, 2 stands for green, 3 stands for blue” scheme, you have to go
 497 through your code and change all the 1s, 2s, 3s, and so on. If you use an
 498 enumerated type, you can continue adding elements to the list just by putting
 499 them into the type definition and recompiling.

500 *Use enumerated types as an alternative to boolean variables*

501 Often, a boolean variable isn’t rich enough to express the meanings it needs to.
 502 For example, suppose you have a routine return *True* if it has successfully
 503 performed its task and *False* otherwise. Later you might find that you really have
 504 two kinds of *False*. The first kind means that the task failed, and the effects are
 505 limited to the routine itself; the second kind means that the task failed, and
 506 caused a fatal error that will need to be propagated to the rest of the program. In
 507 this case, an enumerated type with the values *Status_Success*, *Status_Warning*,
 508 and *Status_FatalError* would be more useful than a boolean with the values *True*
 509 and *False*. This scheme can easily be expanded to handle additional distinctions
 510 in the kinds of success or failure.

511 *Check for invalid values*

512 When you test an enumerated type in an *if* or *case* statement, check for invalid
 513 values. Use the *else* clause in a *case* statement to trap invalid values:

514 **Good Visual Basic Example of Checking for Invalid Values in an** 515 **Enumerated Type**

```
516 Select Case screenColor
517     Case Color_Red
518         ...
519     Case Color_Blue
520         ...
521     Case Color_Green
522         ...
523     Case Else
524         DisplayInternalError( False, "Internal Error 752: Invalid color." )
525 End Select
```

523 Here's the test for the invalid
 524 value.

526 *Define the first and last entries of an enumeration for use as loop limits*

527 Defining the first and last elements in an enumeration to be *Color_First*,
 528 *Color_Last*, *Country_First*, *Country_Last*, and so on allows you to write a loop
 529 that loops through the elements of an enumeration. You set up the enumerated
 530 type using explicit values, as shown below:

Visual Basic Example of Setting *First* and *Last* Values in an Enumerated Type

```
Public Enum Country
    Country_First = 0
    Country_China = 0
    Country_England = 1
    Country_France = 2
    Country_Germany = 3
    Country_India = 4
    Country_Japan = 5
    Country_Usa = 6
    Country_Last = 6
End Enum
```

Now the *Country_First* and *Country_Last* values can be used as loop limits, as shown below:

Good Visual Basic Example of Looping Through Elements in an Enumeration

```
' compute currency conversions from US currency to target currency
Dim usaCurrencyConversionRate( Country_Last ) As Single
Dim iCountry As Country
For iCountry = Country_First To Country_Last
    usaCurrencyConversionRate( iCountry ) = ConversionRate( Country_Usa, iCountry )
Next
```

Reserve the first entry in the enumerated type as invalid

When you declare an enumerated type, reserve the first value as an invalid value. Examples of this were shown earlier in the Visual Basic declarations of *Color*, *Country*, and *Output* types. Many compilers assign the first element in an enumerated type to the value 0. Declaring the element that's mapped to 0 to be invalid helps to catch variables that were not properly initialized since they are more likely to be 0 than any other invalid value.

Here is how the *Country* declaration would look with that approach:

Visual Basic Example of Declaring the First Value in an Enumeration to be Invalid

```
Public Enum Country
    Country_InvalidFirst = 0
    Country_First = 1
    Country_China = 1
    Country_England = 2
    Country_France = 3
    Country_Germany = 4
```



```
571 Country_India = 5
572 Country_Japan = 6
573 Country_Usa = 7
574 Country_Last = 7
575 End Enum
```

Define precisely how First and Last elements are to be used in the project coding standard, and use them consistently

Using *InvalidFirst*, *First*, and *Last* elements in enumerations can make array declarations and loops more readable. But it has the potential to create confusion about whether the valid entries in the enumeration begin at 0 or 1 and whether the first and last elements of the enumeration are valid. If this technique is used, the project's coding standard should require that *InvalidFirst*, *First*, and *Last* elements be used consistently in all enumerations to reduce errors.

Beware of pitfalls of assigning explicit values to elements of an enumeration

Some languages allow you to assign specific values to elements within an enumeration, as shown in the C++ example below:

C++ Example of Explicitly Assigning Values to an Enumeration

```
enum Color {
    Color_InvalidFirst = 0,
    Color_Red = 1,
    Color_Green = 2,
    Color_Blue = 4,
    Color_InvalidLast = 8
};
```

In this C++ example, if you declared a loop index of type *Color* and attempted to loop through *Colors*, you would loop through the invalid values of 3, 5, 6, and 7 as well as the valid values of 1, 2, and 4.

If Your Language Doesn't Have Enumerated Types

If your language doesn't have enumerated types, you can simulate them with global variables or classes. For example, here are declarations you could use in Java:

603 **CROSS-REFERENCE** At
604 the time I'm writing this,
605 Java does not support
606 enumerated types. By the
607 time you read this, it
608 probably will. This is a good
609 example of the "rolling wave
610 of technology" discussed in
611 Section 4.3, "Your Location
on the Technology Wave."

612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634

635
636
637

Java Example of Simulating Enumerated Types

```
// set up Color enumerated type
class Color {
    private Color() {}
    public static final Color Red = new Color();
    public static final Color Green = new Color();
    public static final Color Blue = new Color();
}

// set up Country enumerated type
class Country {
    private Country() {}
    public static final Country China = new Country();
    public static final Country England = new Country();
    public static final Country France = new Country();
    public static final Country Germany = new Country();
    public static final Country India = new Country();
    public static final Country Japan = new Country();
}

// set up Output enumerated type
class Output {
    private Output() {}
    public static final Output Screen = new Output();
    public static final Output Printer = new Output();
    public static final Output File = new Output();
}
```

These enumerated types make your program more readable because you can use the public class members such as *Color.Red* and *Country.England* instead of named constants. This particular method of creating enumerated types is also typesafe; because each type is declared as a class, the compiler will check for invalid assignments such as *Output output = Country.England* (Bloch 2001).

In languages that don't support classes, the same basic effect could be achieved through disciplined use of global variables for each of the elements of the enumeration.

12.7 Named Constants

A named constant is like a variable except that you can't change the constant's value once you've assigned it. Named constants enable you to refer to fixed quantities such as the maximum number of employees by a name rather than a number—*MaximumEmployees* rather than *1000*, for instance.

Using a named constant is a way of “parameterizing” your program—putting an aspect of your program that might change into a parameter that you can change in one place rather than having to make changes throughout the program. If you have ever declared an array to be as big as you think it will ever need to be and then run out of space because it wasn’t big enough, you can appreciate the value of named constants. When an array size changes, you change only the definition of the constant you used to declare the array. This “single-point control” goes a long way toward making software truly “soft”—easy to work with and change.

Use named constants in data declarations

Using named constants helps program readability and maintainability in data declarations and in statements that need to know the size of the data they are working with. In the example below, you use *PhoneLength_c* to describe the length of employee phone numbers rather than the literal 7.

Good Visual Basic Example of Using a Named Constant in a Data Declaration

```

658 Const AREA_CODE_LENGTH = 3
659 LOCAL_NUMBER_LENGTH = 7
660     is declared as a constant
661         here.
662     Type PHONE_NUMBER
663         areaCode( AREA_CODE_LENGTH ) As String
664         localNumber( LOCAL_NUMBER_LENGTH ) As String
665     End Type
666     ...
667     ' make sure all characters in phone number are digits
668     For iDigit = 1 To LOCAL_NUMBER_LENGTH
669         If ( phoneNumber.localNumber( iDigit ) < "0" ) Or _
670             ( "9" < phoneNumber.localNumber( iDigit ) ) Then
671             ' do some error processing
672             ...

```

This is a simple example, but you can probably imagine a program in which the information about the phone-number length is needed in many places.

At the time you create the program, the employees all live in one country, so you need only seven digits for their phone numbers. As the company expands and branches are established in different countries, you’ll need longer phone numbers. If you have parameterized, you can make the change in only one place: in the definition of the named constant *LOCAL_NUMBER_LENGTH*.

FURTHER READING For more details on the value of single-point control, see pages 57-60 of *Software Conflict* (Glass 1991).

As you might expect, the use of named constants has been shown to greatly aid program maintenance. As a general rule, any technique that centralizes control over things that might change is a good technique for reducing maintenance efforts (Glass 1991).

683
684

Avoid literals, even “safe” ones
In the loop below, what do you think the *12* represents?

685

Visual Basic Example of Unclear Code

686

```
For i = 1 To 12  
    profit( i ) = revenue( i ) - expense( i )  
Next
```

687

688

689

690

691

Because of the specific nature of the code, it appears that the code is probably looping through the 12 months in a year. But are you *sure*? Would you bet your Monty Python collection on it?

692

693

694

695

In this case, you don't need to use a named constant to support future flexibility: it's not very likely that the number of months in a year will change anytime soon. But if the way the code is written leaves any shadow of a doubt about its purpose, clarify it with a well-named constant, as shown below.

696

Visual Basic Example of Clearer Code

697

698

699

700

701

```
For i = 1 To NUM_MONTHS_IN_YEAR  
    profit( i ) = revenue( i ) - expense( i )  
Next
```

This is better, but, to complete the example, the loop index should also be named something more informative.

702

Visual Basic Example of Even Clearer Code

703

704

705

706

707

```
For month = 1 To NUM_MONTHS_IN_YEAR  
    profit( month ) = revenue( month ) - expense( month )  
Next
```

This example seems quite good, but we can push it even one step further through using an enumerated type:

708

Visual Basic Example of Very Clear Code

709

710

711

712

```
For month = Month_January To Month_December  
    profit( month ) = revenue( month ) - expense( month )  
Next
```

With this final example, there can be no doubt about the purpose of the loop.

713

714

715

Even if you think a literal is safe, use named constants instead. Be a fanatic about rooting out literals in your code. Use a text editor to search for 2, 3, 4, 5, 6, 7, 8, and 9 to make sure you haven't used them accidentally.

716 **CROSS-REFERENCE** For
717 details on simulating
718 enumerated types, see “If
719 Your Language Doesn’t
720 Have Enumerated Types” in
721 Section 12.6, earlier in this
chapter.

Simulate named constants with appropriately scoped variables or classes

If your language doesn’t support named constants, you can create your own. By using an approach similar to the approach suggested in the earlier Java example in which enumerated types were simulated, you can gain many of the advantages of named constants. Typical scoping rules apply—prefer local scope, class scope, and global scope in that order.

Use named constants consistently

It’s dangerous to use a named constant in one place and a literal in another to represent the same entity. Some programming practices beg for errors; this one is like calling an 800 number and having errors delivered to your door. If the value of the named constant needs to be changed, you’ll change it and think you’ve made all the necessary changes. You’ll overlook the hard-coded literals, your program will develop mysterious defects and fixing them will be a lot harder than picking up the phone and yelling for help.

12.8 Arrays

Arrays are the simplest and most common type of structured data. In some languages, arrays are the only type of structured data. An array contains a group of items that are all of the same type and that are directly accessed through the use of an array index. Here are some tips on using arrays.

KEY POINT

Make sure that all array indexes are within the bounds of the array

In one way or another, all problems with arrays are caused by the fact that array elements can be accessed randomly. The most common problem arises when a program tries to access an array element that’s out of bounds. In some languages, this produces an error; in others, it simply produces bizarre and unexpected results.

Think of arrays as sequential structures

Some of the brightest people in computer science have suggested that arrays never be accessed randomly, but only sequentially (Mills and Linger 1986). Their argument is that random accesses in arrays are similar to random *gotos* in a program: Such accesses tend to be undisciplined, error prone, and hard to prove correct. Instead of arrays, they suggest using sets, stacks, and queues, whose elements are accessed sequentially.

HARD DATA

In a small experiment, Mills and Linger found that designs created this way resulted in fewer variables and fewer variable references. The designs were relatively efficient and led to highly reliable software.

751 Consider using container classes that you access sequentially—sets, stacks,
 752 queues, and so on—as alternatives before you automatically choose an array.

753 **CROSS-REFERENCE** Issu
 754 es in using arrays and loops
 755 are similar and related. For
 756 details on loops, see Chapter
 757 16, “Controlling Loops.”

Check the end points of arrays

Just as it’s helpful to think through the end points in a loop structure, you can catch a lot of errors by checking the end points of arrays. Ask yourself whether the code correctly accesses the first element of the array or mistakenly accesses the element before or after the first element. What about the last element? Will the code make an off-by-one error? Finally, ask yourself whether the code correctly accesses the middle elements of the array.

If an array is multidimensional, make sure its subscripts are used in the correct order

It’s easy to say `Array[i][j]` when you mean `Array[j][i]`, so take the time to double-check that the indexes are in the right order. Consider using more meaningful names than *i* and *j* in cases in which their roles aren’t immediately clear.

Watch out for index cross talk

If you’re using nested loops, it’s easy to write `Array[j]` when you mean `Array[i]`. Switching loop indexes is called “index cross talk.” Check for this problem. Better yet, use more meaningful index names than *i* and *j* and make it harder to commit cross-talk mistakes in the first place.

Throw in an extra element at the end of an array

Off-by-one errors are common with arrays. If your array access is off by one and you write beyond the end of an array, you can cause a serious error. When you declare the array to be one bigger than the size you think you’ll need, you give yourself a cushion and soften the consequences of an off-by-one error.

This is admittedly a sloppy way to program, and you should consider what you’re saying about yourself before you do it. But if you decide that it’s the least of your evils, it can be an effective safeguard.

In C, use the `ARRAY_LENGTH()` macro to work with arrays

You can build extra flexibility into your work with arrays by defining an `ARRAY_LENGTH()` macro that looks like this:

C Example of Defining an `ARRAY_LENGTH()` Macro

```
#define ARRAY_LENGTH( x ) (sizeof(x)/sizeof(x[0]))
```

When you use operations on an array, instead of using a named constant for the upper bound of the array size, use the `ARRAY_LENGTH()` macro. Here’s an example:

787

788

789

790

791

792

793

794

795

796

797

798

799

800

801

802

803

804

805

806

807

808

809

810

811

812

813

814

815

816

817

818

819

820

821

822

823

C Example of Using the `ARRAY_LENGTH()` Macro for Array Operations

```
ConsistencyRatios[] =
    { 0.0, 0.0, 0.58, 0.90, 1.12,
      1.24, 1.32, 1.41, 1.45, 1.49,
      1.51, 1.48, 1.56, 1.57, 1.59 };
    ...
    for ( RatioIdx = 0; RatioIdx < ARRAY_LENGTH( ConsistencyRatios ); RatioIdx++ );
    ...
```

This technique is particularly useful for dimensionless arrays such as the one in the example. If you add or subtract entries, you don't have to remember to change a named constant that describes the array's size. Of course, the technique works with dimensioned arrays too, but if you use this approach, you don't always need to set up an extra named constant for the array definition.

12.9 Creating Your Own Types

Programmer-defined variable types are one of the most powerful capabilities a language can give you to clarify your understanding of a program. They protect your program against unforeseen changes and make it easier to read—all without requiring you to design, construct, and test new classes. If you're using C, C++ or another language that allows user-defined types, take advantage of them!

To appreciate the power of type creation, suppose you're writing a program to convert coordinates in an x, y, z system to latitude, longitude, and elevation. You think that double-precision floating-point numbers might be needed but would prefer to write a program with single-precision floating-point numbers until you're absolutely sure. You can create a new type specifically for coordinates by using a `typedef` statement in C or C++ or the equivalent in another language. Here's how you'd set up the type definition in C++:

C++ Example of Creating a Type

```
typedef float Coordinate; // for coordinate variables
```

This type definition declares a new type, *Coordinate*, that's functionally the same as the type *float*. To use the new type, you declare variables with it just as you would with a predefined type such as *float*. Here's an example:

C++ Example of Using the Type You've Created

```
Routine1( ... ) {
    Coordinate latitude;    // latitude in degrees
    Coordinate longitude;   // longitude in degrees
    Coordinate elevation;   // elevation in meters from earth center
    ...
```

```

824 }
825 ...
826
827 Routine2( ... ) {
828     Coordinate x;    // x coordinate in meters
829     Coordinate y;    // y coordinate in meters
830     Coordinate z;    // z coordinate in meters
831     ...
832 }

```

In this code, the variables *latitude*, *longitude*, *elevation*, *x*, *y*, and *z* are all declared to be of type *Coordinate*.

Now suppose that the program changes and you find that you need to use double-precision variables for coordinates after all. Because you defined a type specifically for coordinate data, all you have to change is the type definition. And you have to change it in only one place: in the *typedef* statement. Here's the changed type definition:

C++ Example of Changed Type Definition

```
typedef double Coordinate; // for coordinate variables
```

Here's a second example—this one in Pascal. Suppose you're creating a payroll system in which employee names are a maximum of 30 characters long. Your users have told you that no one *ever* has a name longer than 30 characters. Do you hard-code the number 30 throughout your program? If you do, you trust your users a lot more than I trust mine! A better approach is to define a type for employee names:

Pascal Example of Creating a Type for Employee Names

```
Type
  EmployeeName_t = array[ 1..30 ] of char;
```

When a string or an array is involved, it's usually wise to define a named constant that indicates the length of the string or array and then use the named constant in the type definition. You'll find many places in your program in which to use the constant—this is just the first place in which you'll use it. Here's how it looks:

Pascal Example of Better Type Creation

```
Const
  NAMELENGTH_C = 30;
  ...
Type
  EmployeeName_t = array[ 1..NAMELENGTH_C ] of char;
```


862 A more powerful example would combine the idea of creating your own types
863 with the idea of information hiding. In some cases, the information you want to
864 hide is information about the type of the data.

865 The coordinates example in C++ is about halfway to information hiding. If you
866 always use *Coordinate* rather than *float* or *double*, you effectively hide the type
867 of the data. In C++, this is about all the information hiding the language does for
868 you. For the rest, you or subsequent users of your code have to have the
869 discipline not to look up the definition of *Coordinate*. C++ gives you figurative,
870 rather than literal, information-hiding ability.

871 Other languages such as Ada go a step further and support literal information
872 hiding. Here's how the *Coordinate* code fragment would look in an Ada package
873 that declares it:

874 Ada Example of Hiding Details of a Type Inside a Package

```
875 package Transformation is  
876     type Coordinate is private;  
877     ...  
878
```

This statement declares
Coordinate as private to the
package.

Here's how *Coordinate* looks in another package, one that uses it:

879 Ada Example of Using a Type from Another Package

```
880 with Transformation;  
881 ...  
882 procedure Routine1(...) ...  
883     latitude: Coordinate;  
884     longitude: Coordinate;  
885 begin  
886     -- statements using latitude and longitude  
887     ...  
888 end Routine1;
```

889 Notice that the *Coordinate* type is declared as *private* in the package
890 specification. That means that the only part of the program that knows the
891 definition of the *Coordinate* type is the private part of the *Transformation*
892 package. In a development environment with a group of programmers, you could
893 distribute only the package specification, which would make it harder for a
894 programmer working on another package to look up the underlying type of
895 *Coordinate*. The information would be literally hidden. Languages like C++ that
896 require you to distribute the definition of *Coordinate* in header files undermine
897 true information hiding.

898 These examples have illustrated several reasons to create your own types:

- 899 • To make modifications easier. It's little work to create a new type, and it
900 gives you a lot of flexibility.
- 901 • To avoid excessive information distribution. Hard typing spreads data-typing
902 details around your program instead of centralizing them in one place. This
903 is an example of the information-hiding principle of centralization discussed
904 in Section 6.2.
- 905 • To increase reliability. In Ada you can define types such as *type Age_t is*
906 *range 0..99*. The compiler then generates run-time checks to verify that any
907 variable of type *Age_t* is always within the range *0..99*.
- 908 • To make up for language weaknesses. If your language doesn't have the
909 predefined type you want, you can create it yourself. For example, C doesn't
910 have a boolean or logical type. This deficiency is easy to compensate for by
911 creating the type yourself:
912

```
typedef int Boolean_t;
```

913 Why Are the Examples of Creating Your Own 914 Types in Pascal and Ada?

915 Pascal and Ada have gone the way of the stegosaurus and, in general, the
916 languages that have replaced them are more usable. In the area of simple type
917 definitions, however, I think C++, Java, and Visual Basic represent a case of
918 three steps forward and one step back. An Ada declaration like

```
919     currentTemperature: INTEGER range 0..212;  
920 contains important semantic information that a statement like
```

```
921     int temperature;  
922 does not. Going a step further, a type declaration like
```

```
923     type Temperature is range 0..212;  
924     ...  
925     currentTemperature: Temperature;
```

926 allows the compiler to ensure that *currentTemperature* is assigned only to other
927 variables with the *Temperature* type, and very little extra coding is required to
928 provide that extra safety margin.

929 Of course a programmer could create a *Temperature* class to enforce the same
930 semantics that were enforced automatically by the Ada language, but the step
931 from creating a simple data type in one line of code to creating a class is a big
932 step. In many situations, a programmer would create the simple type but would
933 not step up to the additional effort of creating a class.

934

935 **CROSS-REFERENCE** In
936 each case, consider whether
937 creating a class might work
938 better than a simple data type.
939 For details, see Chapter 6,
940 “Working Classes.”

941

942

943

944

945

946

947

948

949

950

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

966

967

968

969

970

Guidelines for Creating Your Own Types

Here are a few guidelines to keep in mind as you create your own “user-defined” types:

Create types with functionally oriented names

Avoid type names that refer to the kind of computer data underlying the type. Use type names that refer to the parts of the real-world problem that the new type represents. In the examples above, the definitions created well-named types for coordinates and names—real-world entities. Similarly, you could create types for currency, payment codes, ages, and so on—aspects of real-world problems.

Be wary of creating type names that refer to predefined types. Type names like *BigInteger* or *LongString* refer to computer data rather than the real-world problem. The big advantage of creating your own type is that it provides a layer of insulation between your program and the implementation language. Type names that refer to the underlying programming-language types poke holes in the insulation. They don’t give you much advantage over using a predefined type. Problem-oriented names, on the other hand, buy you easy modifiability and data declarations that are self-documenting.

Avoid predefined types

If there is any possibility that a type might change, avoid using predefined types anywhere but in *typedef* or *type* definitions. It’s easy to create new types that are functionally oriented, and it’s hard to change data in a program that uses hard-wired types. Moreover, use of functionally oriented type declarations partially documents the variables declared with them. A declaration like *Coordinate x* tells you a lot more about *x* than a declaration like *float x*. Use your own types as much as you can.

Don’t redefine a predefined type

Changing the definition of a standard type can create confusion. For example, if your language has a predefined type *Integer*, don’t create your own type called *Integer*. Readers of your code might forget that you’ve redefined the type and assume that the *Integer* they see is the *Integer* they’re used to seeing.

Define substitute types for portability

In contrast to the advice that you not change the definition of a standard type, you might want to define substitutes for the standard types so that on different hardware platforms you can make the variables represent exactly the same entities. For example, you can define a type *INT* and use it instead of *int*, or a type *LONG* instead of *long*. Originally, the only difference between the two types would be their capitalization. But when you moved the program to a new

hardware platform, you could redefine the capitalized versions so that they could match the data types on the original hardware.

If your language isn't case sensitive, you'll have to differentiate the names by some means other than capitalization.

Consider creating a class rather than using a typedef
Simple typedefs can go a long way toward hiding information about a variable's underlying type. In some cases, however, you might want the additional flexibility and control you'll achieve by creating a class. For details, see Chapter 6, "Working Classes."

CHECKLIST: Fundamental Data

Numbers in General

- ☐ Does the code avoid magic numbers?
- ☐ Does the code anticipate divide-by-zero errors?
- ☐ Are type conversions obvious?
- ☐ If variables with two different types are used in the same expression, will the expression be evaluated as you intend it to be?
- ☐ Does the code avoid mixed-type comparisons?
- ☐ Does the program compile with no warnings?

Integers

- ☐ Do expressions that use integer division work the way they're meant to?
- ☐ Do integer expressions avoid integer-overflow problems?

Floating-Point Numbers

- ☐ Does the code avoid additions and subtractions on numbers with greatly different magnitudes?
- ☐ Does the code systematically prevent rounding errors?
- ☐ Does the code avoid comparing floating-point numbers for equality?

Characters and Strings

- ☐ Does the code avoid magic characters and strings?
- ☐ Are references to strings free of off-by-one errors?
- ☐ Does C code treat string pointers and character arrays differently?
- ☐ Does C code follow the convention of declaring strings to be length *constant+1*?
- ☐ Does C code use arrays of characters rather than pointers, when appropriate?

CROSS-REFERENCE For a checklist that applies to general data issues rather than to issues with specific types of data, see the checklist in Chapter 10, "General Issues in Using Variables." For a checklist of considerations in naming varieties, see the checklist in Chapter 11, "The Power of Variable Names."

- 1004 ☐ Does C code initialize strings to *NULLs* to avoid endless strings?
- 1005 ☐ Does C code use *strncpy()* rather than *strcpy()*? And *strncat()* and *strncmp()*?

1006 **Boolean Variables**

- 1007 ☐ Does the program use additional boolean variables to document conditional tests?
- 1008
- 1009 ☐ Does the program use additional boolean variables to simplify conditional tests?
- 1010

1011 **Enumerated Types**

- 1012 ☐ Does the program use enumerated types instead of named constants for their improved readability, reliability, and modifiability?
- 1013
- 1014 ☐ Does the program use enumerated types instead of boolean variables when a variable's use cannot be completely captured with *TRUE* and *FALSE*?
- 1015
- 1016 ☐ Do tests using enumerated types test for invalid values?
- 1017 ☐ Is the first entry in an enumerated type reserved for "invalid"?

1018 **Named Constants**

- 1019 ☐ Does the program use named constants for data declarations and loop limits rather than magic numbers?
- 1020
- 1021 ☐ Have named constants been used consistently—not named constants in some places, literals in others?
- 1022

1023 **Arrays**

- 1024 ☐ Are all array indexes within the bounds of the array?
- 1025 ☐ Are array references free of off-by-one errors?
- 1026 ☐ Are all subscripts on multidimensional arrays in the correct order?
- 1027 ☐ In nested loops, is the correct variable used as the array subscript, avoiding loop-index cross talk?
- 1028

1029 **Creating Types**

- 1030 ☐ Does the program use a different type for each kind of data that might change?
- 1031
- 1032 ☐ Are type names oriented toward the real-world entities the types represent rather than toward programming-language types?
- 1033
- 1034 ☐ Are the type names descriptive enough to help document data declarations?
- 1035 ☐ Have you avoided redefining predefined types?
- 1036 ☐ Have you considered creating a new class rather than simply redefining a type?
- 1037
- 1038
-

Key Points

- Working with specific data types means remembering many individual rules for each type. Use the checklist to make sure that you've considered the common problems.
- Creating your own types makes your programs easier to modify and more self-documenting, if your language supports that capability.
- When you create a simple type using *typedef* or its equivalent, consider whether you should be creating a new class instead.