

17

Unusual Control Structures

CC2E.COM/1778

Contents

17.1 Multiple Returns from a Routine

17.2 Recursion

17.3 *goto*

17.4 Perspective on Unusual Control Structures

Related Topics

General control issues: Chapter 19

Straight-line code: Chapter 14

Code with conditionals: Chapter 15

Code with loops: Chapter 16

Exception handling: Section 8.4

SEVERAL CONTROL CONSTRUCTS exist in a hazy twilight zone somewhere between being leading-edge and being discredited and disproved—often in both places at the same time! These constructs aren't available in all languages but can be useful when used with care in those languages that do offer them.

17.1 Multiple Returns from a Routine

Most languages support some means of exiting from a routine partway through the routine. The *return* and *exit* statements are control constructs that enable a program to exit from a routine at will. They cause the routine to terminate through the normal exit channel, returning control to the calling routine. The word *return* is used here as a generic term for *return* in C++ and Java, *Exit Sub* and *Exit Function* in Visual Basic, and similar constructs. Here are guidelines for using the *return* statement:

KEY POINT

Use a return when it enhances readability

In certain routines, once you know the answer, you want to return it to the calling routine immediately. If the routine is defined in such a way that it doesn't

require any further cleanup once it detects an error, not returning immediately means that you have to write more code.

The following is a good example of a case in which returning from multiple places in a routine makes sense:

C++ Example of a Good Multiple Return from a Routine

```
COMPARISON Compare ( int value1, int value2 ) {  
    if ( value1 < value2 ) {  
        return Comparison_LessThan;  
    }  
    else if ( value1 > value2 ) {  
        return Comparison_GreaterThan;  
    }  
    else {  
        return Comparison_Equal;  
    }  
}
```

Other examples are less clear-cut, as the next section illustrates.

Use guard clauses (early returns or exits) to simplify complex error processing

Code that has to check for numerous error conditions before performing its nominal actions can result in deeply indented code and can obscure the nominal case, as shown here:

Visual Basic Code That Obscures the Nominal Case

```
If file.validName() Then  
    If file.Open() Then  
        If encryptionKey.valid() Then  
            If file.Decrypt( encryptionKey ) Then  
                ' lots of code  
                ...  
            End If  
        End If  
    End If  
End If
```

This is the code for the nominal case.

Indenting the main body of the routine inside four *if* statements is aesthetically ugly, especially if there's much code inside the innermost *if* statement. In such cases, the flow of the code is sometimes clearer if the erroneous cases are checked first, clearing the way for the nominal path through the code. Here's how that might look:

Simple Visual Basic Code That Uses Early Exits to Clarify the Nominal Case

```
' set up, bailing out if errors are found
If Not file.validName() Then Exit Sub
If Not file.Open() Then Exit Sub
If Not encryptionKey.valid() Then Exit Sub
If Not file.Decrypt( encryptionKey ) Then Exit Sub

' lots of code
...
```

The simple code above makes this technique look like a tidy solution, but production code often requires more extensive housekeeping or cleanup when an error condition is detected. Here is a more realistic example:

More Realistic Visual Basic Code That Uses Early Exits to Clarify the Nominal Case

```
' set up, bailing out if errors are found
If Not file.validName() Then
    errorStatus = FileError_InvalidFileName
    Exit Sub
End If

If Not file.Open() Then
    errorStatus = FileError_CantOpenFile
    Exit Sub
End If

If Not encryptionKey.valid() Then
    errorStatus = FileError_InvalidEncryptionKey
    Exit Sub
End If

If Not file.Decrypt( encryptionKey ) Then
    errorStatus = FileError_CantDecryptFile
    Exit Sub
End If

' lots of code
...
```

This is the code for the nominal case.

With production-size code, the *Exit Sub* approach creates a noticeable amount of code before the nominal case is handled. The *Exit Sub* approach does avoid the deep nesting of the first example, however, and, if the code in the first example were expanded to show setting an *errorStatus* variable, the *Exit Sub* approach would do a better job of keeping related statements together. When all the dust

settles, the *Exit Sub* approach does appear more readable and maintainable, just not by a very wide margin.

Minimize the number of returns in each routine

It's harder to understand a routine if, reading it at the bottom, you're unaware of the possibility that it returned somewhere above. For that reason, use returns judiciously—only when they improve readability.

17.2 Recursion

In recursion, a routine solves a small part of a problem itself, divides the problem into smaller pieces, and then calls itself to solve each of the smaller pieces. Recursion is usually called into play when a small part of the problem is easy to solve and a large part is easy to decompose into smaller pieces.

KEY POINT

Recursion isn't useful very often, but when used judiciously it produces exceptionally elegant solutions. Here's an example in which a sorting algorithm makes excellent use of recursion:

Java Example of a Sorting Algorithm That Uses Recursion

```
void QuickSort( int firstIndex, int lastIndex, String [] names ) {  
    if ( lastIndex > firstIndex ) {  
        int midPoint = Partition( firstIndex, lastIndex, names );  
        QuickSort( firstIndex, midPoint-1, names );  
        QuickSort( midPoint+1, lastIndex, names );  
    }  
}
```

Here are the recursive calls.


In this case, the sorting algorithm chops an array in two and then calls itself to sort each half of the array. When it calls itself with a subarray that's too small to sort (*lastIndex* <= *firstIndex*), it stops calling itself.

In general, recursion leads to small code and slow execution and chews up stack space. For a small group of problems, recursion can produce simple, elegant solutions. For a slightly larger group of problems, it can produce simple, elegant, hard-to-understand solutions. For most problems, it produces massively complicated solutions—in those cases, simple iteration is usually more understandable. Use recursion selectively.

Example of Recursion

Suppose you have a data type that represents a maze. A maze is basically a grid, and at each point on the grid you might be able to turn left, turn right, move up, or move down. You'll often be able to move in more than one direction.

145 How do you write a program to find its way through the maze? If you use
 146 recursion, the answer is fairly straightforward. You start at the beginning and
 147 then try all possible paths until you find your way out of the maze. The first time
 148 you visit a point, you try to move left. If you can't move left, you try to go up or
 149 down, and if you can't go up or down, you try to go right. You don't have to
 150 worry about getting lost because you drop a few bread crumbs on each spot as
 151 you visit it, and you don't visit the same spot twice.



152

153

154

155

156

157

F17xx01

Figure 17-1

Recursion can be a valuable tool in the battle against complexity—when used to attack suitable problems.

Here's how the recursive code looks:

C++ Example of Moving Through a Maze Recursively

```
bool FindPathThroughMaze( Maze maze, Point position ) {
    // if the position has already been tried, don't try it again
    if ( AlreadyTried( maze, position ) ) {
        return false;
    }

    // if this position is the exit, declare success
    if ( ThisIsTheExit( maze, position ) ) {
        return true;
    }

    // remember that this position has been tried
    RememberPosition( maze, position );

    // check the paths to the left, up, down, and to the right; if
    // any path is successful, stop looking
    if ( MoveLeft( maze, position, &newPosition ) ) {
```

```
176         if ( FindPathThroughMaze( maze, newPosition ) ) {
177             return true;
178         }
179     }
180
181     if ( MoveUp( maze, position, &newPosition ) ) {
182         if ( FindPathThroughMaze( maze, newPosition ) ) {
183             return true;
184         }
185     }
186
187     if ( MoveDown( maze, position, &newPosition ) ) {
188         if ( FindPathThroughMaze( maze, newPosition ) ) {
189             return true;
190         }
191     }
192
193     if ( MoveRight( maze, position, &newPosition ) ) {
194         if ( FindPathThroughMaze( maze, newPosition ) ) {
195             return true;
196         }
197     }
198     return false;
199 }
```

200 The first line of code checks to see whether the position has already been tried.
201 One key aim in writing a recursive routine is the prevention of infinite recursion.
202 In this case, if you don't check for having tried a point, you might keep trying it
203 infinitely.

204 The second statement checks to see whether the position is the exit from the
205 maze. If *ThisIsTheExit()* returns *true*, the routine itself returns *true*.

206 The third statement remembers that the position has been visited. This prevents
207 the infinite recursion that would result from a circular path.

208 The remaining lines in the routine try to find a path to the left, up, down, and to
209 the right. The code stops the recursion if the routine ever returns *true*, that is,
210 when the routine finds a path through the maze.

211 The logic used in this routine is fairly straightforward. Most people experience
212 some initial discomfort using recursion because it's self-referential. In this case,
213 however, an alternative solution would be much more complicated and recursion
214 works well.

Tips for Using Recursion

Here are some tips for using recursion:

Make sure the recursion stops

Check the routine to make sure that it includes a nonrecursive path. That usually means that the routine has a test that stops further recursion when it's not needed. In the maze example, the tests for *AlreadyTried()* and *ThisIsTheExit()* ensure that the recursion stops.

Use safety counters to prevent infinite recursion

If you're using recursion in a situation that doesn't allow a simple test such as the one just described, use a safety counter to prevent infinite recursion. The safety counter has to be a variable that's not re-created each time you call the routine. Use a class member variable or pass the safety counter as a parameter. Here's an example:

Visual Basic Example of Using a Safety Counter to Prevent Infinite Recursion

```
Public Sub RecursiveProc( ByRef safetyCounter As Integer )
    If ( safetyCounter > SAFETY_LIMIT ) Then
        Exit Sub
    End If
    safetyCounter = safetyCounter + 1
    ...
    RecursiveProc( safetyCounter )
End Sub
```

In this case, if the routine exceeds the safety limit, it stops recursing.

If you don't want to pass the safety counter as an explicit parameter, you could use a *static* variable in C++, Java, or Visual Basic, or the equivalent in other languages.

Limit recursion to one routine

Cyclic recursion (A calls B calls C calls A) is dangerous because it's hard to detect. Mentally managing recursion in one routine is tough enough; understanding recursion that spans routines is too much. If you have cyclic recursion, you can usually redesign the routines so that the recursion is restricted to a single routine. If you can't and you still think that recursion is the best approach, use safety counters as a recursive insurance policy.

Keep an eye on the stack

With recursion, you have no guarantees about how much stack space your program uses and it's hard to predict in advance how the program will behave at

The recursive routine must be able to change the value of safetyCounter, so in Visual Basic it's a ByRef parameter.

run time. You can take a couple of steps to control its run-time behavior, however.

First, if you use a safety counter, one of the considerations in setting a limit for it should be how much stack you're willing to allocate to the recursive routine. Set the safety limit low enough to prevent a stack overflow.

Second, watch for allocation of local variables in recursive functions, especially memory-intensive objects. In other words, use *new* to create objects on the heap rather than letting the compiler create *auto* objects on the stack.

Don't use recursion for factorials or Fibonacci numbers

One problem with computer-science textbooks is that they present silly examples of recursion. The typical examples are computing a factorial or computing a Fibonacci sequence. Recursion is a powerful tool, and it's really dumb to use it in either of those cases. If a programmer who worked for me used recursion to compute a factorial, I'd hire someone else. Here's the recursive version of the factorial routine:

CODING HORROR

Java Example of an Inappropriate Solution: Using Recursion to Compute a Factorial

```
int Factorial( int number ) {  
    if ( number == 1 ) {  
        return 1;  
    }  
    else {  
        return number * Factorial( number - 1 );  
    }  
}
```

In addition to being slow and making the use of run-time memory unpredictable, the recursive version of this routine is harder to understand than the iterative version. Here's the iterative version:

Java Example of an Appropriate Solution: Using Iteration to Compute a Factorial

```
int Factorial( int number ) {  
    int intermediateResult = 1;  
    for ( int factor = 2; factor <= number; factor++ ) {  
        intermediateResult = intermediateResult * factor;  
    }  
    return intermediateResult;  
}
```

You can draw three lessons from this example. First, computer-science textbooks aren't doing the world any favors with their examples of recursion. Second, and

291 more important, recursion is a much more powerful tool than its confusing use in
292 computing factorials or Fibonacci numbers would suggest. Third, and most
293 important, you should consider alternatives to recursion before using it. You can
294 do anything with stacks and iteration that you can do with recursion. Sometimes
295 one approach works better; sometimes the other does. Consider both before you
296 choose either one.

297 17.3 *goto*

298 CC2E.COM/1785 You might think the debate related to *gotos* is extinct, but a quick trip through
299 modern source-code repositories like *SourceForge.net* shows that the *goto* is still
300 alive and well and living deep in your company's server. Moreover, modern
301 equivalents of the *goto* debate still crop up in various guises including debates
302 about multiple returns, multiple loop exits, named loop exits, error processing,
303 and exception handling.

304 Here's a summary of the points on each side of the *goto* debate.

305 The Argument Against *gotos*

306 The general argument against *gotos* is that code without *gotos* is higher-quality
307 code. The famous letter that sparked the original controversy was Edsger
308 Dijkstra's "Go To Statement Considered Harmful" in the March 1968
309 *Communications of the ACM*. Dijkstra observed that the quality of code was
310 inversely proportional to the number of *gotos* the programmer used. In
311 subsequent work, Dijkstra has argued that code that doesn't contain *gotos* can
312 more easily be proven correct.

313 Code containing *gotos* is hard to format. Indentation should be used to show
314 logical structure, and *gotos* have an effect on logical structure. Using indentation
315 to show the logical structure of a *goto* and its target, however, is difficult or
316 impossible.

317 Use of *gotos* defeats compiler optimizations. Some optimizations depend on a
318 program's flow of control residing within a few statements. An unconditional
319 *goto* makes the flow harder to analyze and reduces the ability of the compiler to
320 optimize the code. Thus, even if introducing a *goto* produces an efficiency at the
321 source-language level, it may well reduce overall efficiency by thwarting
322 compiler optimizations.

323 Proponents of *gotos* sometimes argue that they make code faster or smaller. But
324 code containing *gotos* is rarely the fastest or smallest possible. Donald Knuth's
325 marvelous, classic article "Structured Programming with go to Statements" gives

several examples of cases in which using *gotos* makes for slower and larger code (Knuth 1974).

In practice, the use of *gotos* leads to the violation of the principle that code should flow strictly from top to bottom. Even if *gotos* aren't confusing when used carefully, once *gotos* are introduced, they spread through the code like termites through a rotting house. If any *gotos* are allowed, the bad creep in with the good, so it's better not to allow any of them.

Overall, experience in the two decades that followed the publication of Dijkstra's letter showed the folly of producing *goto*-laden code. In a survey of the literature, Ben Shneiderman concluded that the evidence supports Dijkstra's view that we're better off without the *goto* (1980), and many modern languages including Java don't even have *gotos*.

The Argument for *gotos*

The argument for the *goto* is characterized by an advocacy of its careful use in specific circumstances rather than its indiscriminate use. Most arguments against *gotos* speak against indiscriminate use. The *goto* controversy erupted when Fortran was the most popular language. Fortran had no presentable loop structures, and in the absence of good advice on programming loops with *gotos*, programmers wrote a lot of spaghetti code. Such code was undoubtedly correlated with the production of low-quality programs but has little to do with the careful use of a *goto* to make up for a gap in a modern language's capabilities.

A well-placed *goto* can eliminate the need for duplicate code. Duplicate code leads to problems if the two sets of code are modified differently. Duplicate code increases the size of source and executable files. The bad effects of the *goto* are outweighed in such a case by the risks of duplicate code.

The *goto* is useful in a routine that allocates resources, performs operations on those resources, and then deallocates the resources. With a *goto*, you can clean up in one section of code. The *goto* reduces the likelihood of your forgetting to deallocate the resources in each place you detect an error.

In some cases, the *goto* can result in faster and smaller code. Knuth's 1974 article cited a few cases in which the *goto* produced a legitimate gain.

Good programming doesn't mean eliminating *gotos*. Methodical decomposition, refinement, and selection of control structures automatically lead to *goto*-free programs in most cases. Achieving *goto*-less code is not the aim, but the outcome, and putting the focus on avoiding *gotos* isn't helpful.

CROSS-REFERENCE For details on using *gotos* in code that allocates resources, see "Error Processing and *gotos*" in this section. See also the discussion of exception handling in Section 8.4, "Exceptions."

362 *The evidence suggests*
 363 *only that deliberately*
 364 *chaotic control structure*
 365 *degrades [programmer]*
 366 *performance. These*
 367 *experiments provide*
 368 *virtually no evidence for*
 369 *the beneficial effect of*
 370 *any specific method of*
 371 *structuring control flow.*
 372 *— B. A. Sheil*
 373

Decades' worth of research with *gotos* failed to demonstrate their harmfulness. In a survey of the literature, B. A. Sheil concluded that unrealistic test conditions, poor data analysis, and inconclusive results failed to support the claim of Shneiderman and others that the number of bugs in code was proportional to the number of *gotos* (1981). Sheil didn't go so far as to conclude that using *gotos* is a good idea—rather that experimental evidence against them was not conclusive.

Finally, the *goto* has been incorporated into many modern languages including Visual Basic, C++ and the Ada language—the most carefully engineered programming language in history. Ada was developed long after the arguments on both sides of the *goto* debate had been fully developed, and after considering all sides of the issue, Ada's engineers decided to include the *goto*.

The Phony *goto* Debate

A primary feature of most *goto* discussions is a shallow approach to the question. The arguer on the “*gotos* are evil” side presents a trivial code fragment that uses *gotos* and then shows how easy it is to rewrite the fragment without *gotos*. This proves mainly that it's easy to write trivial code without *gotos*.

The arguer on the “I can't live without *gotos*” side usually presents a case in which eliminating a *goto* results in an extra comparison or the duplication of a line of code. This proves mainly that there's a case in which using a *goto* results in one less comparison—not a significant gain on today's computers.

Most textbooks don't help. They provide a trivial example of rewriting some code without a *goto* as if that covered the subject. Here's a disguised example of a trivial piece of code from such a textbook:

C++ Example of Code That's Supposed to Be Easy to Rewrite Without *gotos*

```
do {
    GetData( inputFile, data );
    if ( eof( inputFile ) ) {
        goto LOOP_EXIT;
    }
    DoSomething( data );
} while ( data != -1 );
LOOP_EXIT:
```

The book quickly replaces this code with *gotoless* code:

C++ Example of Supposedly Equivalent Code, Rewritten Without *gotos*

```
GetData( inputFile, data );
```

```
while ( ( !eof( inputFile ) ) && ( ( data != -1 ) ) ) do {  
    DoSomething( data );  
    GetData( inputFile, data )  
}
```

This so-called “trivial” example contains an error. In the case in which *data* equals *-1* entering the loop, the translated code detects the *-1* and exits the loop before executing *DoSomething()*. The original code executes *DoSomething()* before the *-1* is detected. The programming book trying to show how easy it is to code without *gotos* translated its own example incorrectly. But the author of that book shouldn’t feel too bad; other books make similar mistakes. Even the pros have difficulty achieving *gotoless* nirvana.

Here’s a faithful translation of the code with no *gotos*:

C++ Example of Truly Equivalent Code, Rewritten Without *gotos*

```
do {  
    GetData( inputFile, data );  
    if ( !eof( inputFile ) ) {  
        DoSomething( data );  
    }  
} while ( ( data != -1 ) && ( !eof( InputFile ) ) );
```

Even with a correct translation of the code, the example is still phony because it shows a trivial use of the *goto*. Such cases are not the ones for which thoughtful programmers choose a *goto* as their preferred form of control.

It would be hard at this late date to add anything worthwhile to the theoretical *goto* debate. What’s not usually addressed, however, is the situation in which a programmer fully aware of the *gotoless* alternatives chooses to use a *goto* to enhance readability and maintainability.

The following sections present cases in which some experienced programmers have argued for using *gotos*. The discussions provide examples of code with *gotos* and code rewritten without *gotos* and evaluate the trade-offs between the versions.

Error Processing and *gotos*

Writing highly interactive code calls for paying a lot of attention to error processing and cleaning up resources when errors occur. Here’s a code example that purges a group of files. The routine first gets a group of files to be purged, and then it finds each file, opens it, overwrites it, and erases it. The routine checks for errors at each step:

Visual Basic Code with *gotos* That Processes Errors and Cleans Up Resources

```
' This routine purges a group of files.
Sub PurgeFiles( ByRef errorState As Error_Code )
    Dim fileIndex As Integer
    Dim fileToPurge As Data_File
    Dim fileList As File_List
    Dim numFilesToPurge As Integer

    MakePurgeFileList( fileList, numFilesToPurge )

    errorState = FileStatus_Success
    fileIndex = 0
    While ( fileIndex < numFilesToPurge )
        fileIndex = fileIndex + 1
        If Not ( FindFile( fileList( fileIndex ), fileToPurge ) ) Then
            errorState = FileStatus_FileFindError
            GoTo END_PROC
        End If

        If Not OpenFile( fileToPurge ) Then
            errorState = FileStatus_FileOpenError
            GoTo END_PROC
        End If

        If Not OverwriteFile( fileToPurge ) Then
            errorState = FileStatus_FileOverwriteError
            GoTo END_PROC
        End If

        if Erase( fileToPurge ) Then
            errorState = FileStatus_FileEraseError
            GoTo END_PROC
        End If

    Wend

    END_PROC:
    DeletePurgeFileList( fileList, numFilesToPurge )
End Sub
```

This routine is typical of circumstances in which experienced programmers decide to use a *goto*. Similar cases come up when a routine needs to allocate and clean up resources like database connections, memory, or temporary files. The alternative to *gotos* in those cases is usually duplicating code to clean up the resources. In such cases, a programmer might balance the evil of the *goto* against

480 the headache of duplicate-code maintenance and decide that the *goto* is the lesser
 481 evil.

482 You can rewrite the routine above in a couple of ways that avoid *gotos*, and both
 483 ways involve trade-offs. Here are the possible rewrite strategies:

484 ***Rewrite with nested if statements***

485 To rewrite with nested *if* statements, nest the *if* statements so that each is
 486 executed only if the previous test succeeds. This is the standard, textbook
 487 programming approach to eliminating *gotos*. Here's a rewrite of the routine
 488 using the standard approach:

489 **CROSS-REFERENCE** C++
 490 programmers might point out
 491 that this routine could easily
 492 be rewritten with *break* and
 493 no *gotos*. For details, see
 494 "Exiting Loops Early" in
 495 Section 16.2.

496
 497
 498
 499
 500
 501 *The While test has been*
 502 *changed to add a test for*
 503 *errorState.*

504
 505
 506
 507
 508
 509
 510
 511
 512
 513
 514
 515
 516
 517
 518 *This line is 13 lines away from*
 519 *the If statement that invokes*
 520 *it.*

521
 522

Visual Basic Code That Avoids GoTos by Using Nested ifs

```
' This routine purges a group of files.
Sub PurgeFiles( ByRef errorState As Error_Code )
    Dim fileIndex As Integer
    Dim fileToPurge As Data_File
    Dim fileList As File_List
    Dim numFilesToPurge As Integer

    MakePurgeFileList( fileList, numFilesToPurge )

    errorState = FileStatus_Success
    fileIndex = 0
    While ( fileIndex < numFilesToPurge And errorState = FileStatus_Success )

        fileIndex = fileIndex + 1

        If FindFile( fileList( fileIndex ), fileToPurge ) Then
            If OpenFile( fileToPurge ) Then
                If OverwriteFile( fileToPurge ) Then
                    If Not Erase( fileToPurge ) Then
                        errorState = FileStatus_FileEraseError
                    End If
                Else ' couldn't overwrite file
                    errorState = FileStatus_FileOverwriteError
                End If
            Else ' couldn't open file
                errorState = FileStatus_FileOpenError
            End If
        Else ' couldn't find file
            errorState = FileStatus_FileFindError
        End If
    Wend
    DeletePurgeFileList( fileList, numFilesToPurge )
End Sub
```

523
524
525

For people used to programming without *gotos*, this code might be easier to read than the *goto* version, and If you use it, you won't have to face an inquisition from the *goto* goon squad.

526 **CROSS-REFERENCE** For
527 more details on indentation
528 and other coding layout
529 issues, see Chapter 31,
530 "Layout and Style." For
531 details on nesting levels, see
532 Section 19.4, "Taming
533 Dangerously Deep Nesting."

The main disadvantage of this nested-*If* approach is that the nesting level is deep. Very deep. To understand the code, you have to keep the whole set of nested *ifs* in your mind at once. Moreover, the distance between the error-processing code and the code that invokes it is too great: The code that sets *errorState* to *FileStatus_FileFindError*, for example, is 13 lines from the *If* statement that invokes it.

532
533
534
535
536

With the *goto* version, no statement is more than 4 lines from the condition that invokes it. And you don't have to keep the whole structure in your mind at once. You can essentially ignore any preceding conditions that were successful and focus on the next operation. In this case, the *goto* version is more readable and more maintainable than the nested-*If* version.

537
538
539
540

Rewrite with a status variable

To rewrite with a status variable (also called a state variable), create a variable that indicates whether the routine is in an error state. In this case, the routine already uses the *errorState* status variable, so you can use that.

541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563

*The While test has been
changed to add a test for
errorState.*

The status variable is tested.

Visual Basic Code That Avoids *gotos* by Using a Status Variable

```
' This routine purges a group of files.
Sub PurgeFiles( ByRef errorState As Error_Code )
    Dim fileIndex As Integer
    Dim fileToPurge As Data_File
    Dim fileList As File_List
    Dim numFilesToPurge As Integer

    MakePurgeFileList( fileList, numFilesToPurge )

    errorState = FileStatus_Success
    fileIndex = 0
    While ( fileIndex < numFilesToPurge ) And ( errorState = FileStatus_Success )

        fileIndex = fileIndex + 1

        If Not FindFile( fileList( fileIndex ), fileToPurge ) Then
            errorState = FileStatus_FileFindError
        End If

        If ( errorState = FileStatus_Success ) Then
            If Not OpenFile( fileToPurge ) Then
                errorState = FileStatus_FileOpenError
            End If
        End If
    End While
End Sub
```

```

564         End If
565     End If
566
567     The status variable is tested. If ( errorState = FileStatus_Success ) Then
568         If Not OverwriteFile( fileToPurge ) Then
569             errorState = FileStatus_FileOverwriteError
570         End If
571     End If
572
573     The status variable is tested. If ( errorState = FileStatus_Success ) Then
574         If Not Erase( fileToPurge ) Then
575             errorState = FileStatus_FileEraseError
576         End If
577     End If
578 Wend
579 DeletePurgeFileList( fileList, numFilesToPurge )
580 End Sub

```

581 The advantage of the status-variable approach is that it avoids the deeply nested
582 *if-then-else* structures of the first rewrite and is thus easier to understand. It also
583 places the action following the *if-then-else* test closer to the test than the nested-
584 *if* approach did and completely avoids *else* clauses.

585 Understanding the nested-*if* version requires some mental gymnastics. The
586 status-variable version is easier to understand because it closely models the way
587 people think about the problem. You find the file. If everything is OK, you open
588 the file. If everything is still OK, you overwrite the file. If everything is still
589 OK,...

590 The disadvantage of this approach is that using status variables isn't as common
591 a practice as it should be. Document their use fully, or some programmers might
592 not understand what you're up to. In this example, the use of well-named
593 enumerated types helps significantly.

594 **Rewrite with try-finally**

595 Some languages, including Visual Basic and Java, provide a *try-finally* statement
596 that can be used to clean up resources under error conditions.

597 To rewrite using the *try-finally* approach, enclose the code that would otherwise
598 need to check for errors inside a *try* block, and place the cleanup code inside a
599 *finally* block. The *try* block specifies the scope of the exception handling, and the
600 *finally* block performs any resource cleanup. The *finally* block will always be
601 called regardless of whether an exception is thrown and regardless of whether
602 the *PurgeFiles()* routine *Catches* any exception that's thrown.

Visual Basic Code That Avoids *gotos* by Using Try-Finally

```
' This routine purges a group of files. Exceptions are passed to the caller.
Sub PurgeFiles()
    Dim fileIndex As Integer
    Dim fileToPurge As Data_File
    Dim fileList As File_List
    Dim numFilesToPurge As Integer
    MakePurgeFileList( fileList, numFilesToPurge )
    Try
        fileIndex = 0
        While ( fileIndex < numFilesToPurge )
            fileIndex = fileIndex + 1
            FindFile( fileList( fileIndex ), fileToPurge )
            OpenFile( fileToPurge )
            OverwriteFile( fileToPurge )
            Erase( fileToPurge )
        Wend
    Finally
        DeletePurgeFileList( fileList, numFilesToPurge )
    End Try
End Sub
```

This approach assumes that all function calls throw exceptions for failures rather than returning error codes.

The advantage of the *try-finally* approach is it achieves the visual simplicity of the *goto* approach without the use of *gotos*. It also avoids the deeply nested *if-then-else* structures.

The limitation of the *try-finally* approach is that it must be implemented consistently throughout a code base. If the code above was part of a code base that used both error codes and exceptions, the code would be required to set an error code for each possible error, and that requirement would make the code above about as complicated as the other approaches. In that context, the *try-finally* structure wouldn't be decisively more attractive than the other approaches.

A final limitation of this approach is that the *try-finally* statement is not available in all languages.

Comparison of the Approaches

Each of the four methods has something to be said for it. The *goto* approach avoids deep nesting and unnecessary tests but of course has *gotos*. The nested-*if* approach avoids *gotos* but is deeply nested and gives an exaggerated picture of the logical complexity of the routine. The status-variable approach avoids *gotos*

and deep nesting but introduces extra tests. The *try-finally* approach avoids both *gotos* and deep nesting, but isn't available in all languages.

The *try-finally* approach is the most straightforward in languages that provide *try-finally* and in code bases that haven't already standardized on another approach. If *try-finally* isn't an option, the status-variable approach is slightly preferable to the first two because it's more readable and it models the problem better, but that doesn't make it the best approach in all circumstances.

Any of these techniques works well when applied consistently to all the code in a project. Consider all the trade-offs, and then make a project-wide decision about which method to favor.

***gotos* and Sharing Code in an *else* Clause**

One challenging situation in which some programmers would use a *goto* is the case in which you have two conditional tests and an *else* clause and want to execute code in one of the conditions and in the *else* clause. Here's an example of a case that could drive someone to *goto*:

CODING HORROR

C++ Example of Sharing Code in an *else* Clause with a *goto*

```
if ( statusOk ) {  
    if ( dataAvailable ) {  
        importantVariable = x;  
        goto MID_LOOP;  
    }  
}  
else {  
    importantVariable = GetValue();  
  
    MID_LOOP:  
  
    // lots of code  
    ...  
}
```

This is a good example because it's logically tortuous—it's nearly impossible to read as it stands, and it's hard to rewrite correctly without a *goto*. If you think you can easily rewrite it without *gotos*, ask someone to review your code! Several expert programmers have rewritten it incorrectly.

You can rewrite the code in several ways. You can duplicate code, put the common code into a routine and call it from two places, or retest the conditions. In most languages, the rewrite will be a tiny bit larger and slower than the original, but it will be extremely close. Unless the code is in a really hot loop, rewrite it without thinking about efficiency.

The best rewrite would be to put the *// lots of code* part into its own routine. Then you can call the routine from the places you would otherwise have used as origins or destinations of *gotos* and preserve the original structure of the conditional. Here’s how it looks:

C++ Example of Sharing Code in an *else* Clause by Putting Common Code into a Routine

```
if ( statusOk ) {
    if ( dataAvailable ) {
        importantVariable = x;
        DoLotsOfCode( importantVariable );
    }
}
else {
    importantVariable = GetValue();
    DoLotsOfCode( importantVariable );
}
```

Normally, writing a new routine is the best approach. Sometimes, however, it’s not practical to put duplicated code into its own routine. In this case you can work around the impractical solution by restructuring the conditional so that you keep the code in the same routine rather than putting it into a new routine. Here’s how it looks:

C++ Example of Sharing Code in an *else* Clause Without a *goto*

```
if ( ( statusOk && dataAvailable ) || !statusOk ) {
    if ( statusOk && dataAvailable ) {
        importantVariable = x;
    }
    else {
        importantVariable = GetValue();
    }

    // lots of code
    ...
}
```

This is a faithful and mechanical translation of the logic in the *goto* version. It tests *statusOK* two extra times and *dataAvailable* one, but the code is equivalent. If retesting the conditionals bothers you, notice that the value of *statusOK* doesn’t need to be tested twice in the first *if* test. You can also drop the test for *dataAvailable* in the second *if* test.

Summary of Guidelines for Using *gotos*

KEY POINT

721 Use of *gotos* is a matter of religion. My dogma is that in modern languages, you
722 can easily replace nine out of ten *gotos* with equivalent sequential constructs. In
723 these simple cases, you should replace *gotos* out of habit. In the hard cases, you
724 can still exorcise the *goto* in nine out of ten cases: You can break the code into
725 smaller routines, use nested *ifs*, test and retest a status variable, or restructure a
726 conditional. Eliminating the *goto* is harder in these cases, but it's good mental
727 exercise and the techniques discussed in this section give you the tools to do it.

728 In the remaining one case out of 100 in which a *goto* is a legitimate solution to
729 the problem, document it clearly and use it. If you have your rain boots on, it's
730 not worth walking around the block to avoid a mud puddle. But keep your mind
731 open to *gotoless* approaches suggested by other programmers. They might see
732 something you don't.

733 Here's a summary of guidelines for using *gotos*:

- 734 ● Use *gotos* to emulate structured control constructs in languages that don't
735 support them directly. When you do, emulate them exactly. Don't abuse the
736 extra flexibility the *goto* gives you.
- 737 ● Don't use the *goto* when an equivalent built-in construct is available.
- 738 **CROSS-REFERENCE** For
739 details on improving
740 efficiency, see Chapter 25,
741 "Code-Tuning Strategies,"
742 and Chapter 26, "Code-
743 Tuning Techniques."
- 744 ● Measure the performance of any *goto* used to improve efficiency. In most
745 cases, you can recode without *gotos* for improved readability and no loss in
746 efficiency. If your case is the exception, document the efficiency
747 improvement so that *gotoless* evangelists won't remove the *goto* when they
748 see it.
- 749 ● Limit yourself to one *goto* label per routine unless you're emulating
750 structured constructs.
- 751 ● Limit yourself to *gotos* that go forward, not backward, unless you're
752 emulating structured constructs.
- 753 ● Make sure all *goto* labels are used. Unused labels might be an indication of
missing code, namely the code that goes to the labels. If the labels aren't
used, delete them.
- Make sure a *goto* doesn't create unreachable code.
- If you're a manager, adopt the perspective that a battle over a single *goto*
isn't worth the loss of the war. If the programmer is aware of the alternatives
and is willing to argue, the *goto* is probably OK.

17.4 Perspective on Unusual Control Structures

At one time or another, someone thought that each of the following control structures was a good idea:

- Unrestricted use of *gotos*
- Ability to compute a *goto* target dynamically, and jump to the computed location
- Ability to use *goto* to jump from the middle of one routine into the middle of another routine
- Ability to call a routine with a line number or label that allowed execution to begin somewhere in the middle of the routine
- Ability to have the program generate code on the fly, then execute the code it just wrote

At one time, each of these ideas was regarded as acceptable or even desirable, even though now they all look hopelessly quaint, outdated or dangerous. The field of software development has advanced largely through *restricting* what programmers can do with their code. Consequently, I view unconventional control structures with strong skepticism. I suspect that the majority of constructs in this chapter will eventually find their way onto the programmer's scrap heap along with computed *goto* labels, variable routine entry points, self-modifying code, and other structures that favored flexibility and convenience over structure and ability to manage complexity.

CC2E.COM/1792

Additional Resources

Returns

Fowler, Martin. *Refactoring: Improving the Design of Existing Code*, Reading, Mass.: Addison Wesley, 1999. In the description of the refactoring called "Replace Nested Conditional with Guard Clauses," Fowler suggests using multiple *return* statements from a routine to reduce nesting in a set of *if* statements. Fowler argues that multiple *returns* are an appropriate means of achieving greater clarity, and that no harm arises from having multiple returns from a routine.

785

gotos

786

These articles contain the whole *goto* debate. It erupts from time to time in most workplaces, textbooks, and magazines, but you won't hear anything that wasn't fully explored 20 years ago.

787

788

789 CC2E.COM/1799

Dijkstra, Edsger. "Go To Statement Considered Harmful." *Communications of the ACM* 11, no. 3 (March 1968): 147–48, also available from www.cs.utexas.edu/users/EWD/. This is the famous letter in which Dijkstra put the match to the paper and ignited one of the longest-running controversies in software development.

790

791

792

793

794

Wulf, W. A. "A Case Against the GOTO." *Proceedings of the 25th National ACM Conference*, August 1972: 791–97. This paper was another argument against the indiscriminate use of *gotos*. Wulf argued that if programming languages provided adequate control structures, *gotos* would become largely unnecessary. Since 1972, when the paper was written, languages such as C++, Java, and Visual Basic have proven Wulf correct.

795

796

797

798

799

800

Knuth, Donald. "Structured Programming with go to Statements," 1974. In *Classics in Software Engineering*, edited by Edward Yourdon. Englewood Cliffs, N. J.: Yourdon Press, 1979. This long paper isn't entirely about *gotos*, but it includes a horde of code examples that are made more efficient by eliminating *gotos* and another horde of code examples that are made more efficient by adding *gotos*.

801

802

803

804

805

806

Rubin, Frank. "'GOTO Considered Harmful' Considered Harmful." *Communications of the ACM* 30, no. 3 (March 1987): 195–96. In this rather hotheaded letter to the editor, Rubin asserts that *gotoless* programming has cost businesses "hundreds of millions of dollars." He then offers a short code fragment that uses a *goto* and argues that it's superior to *gotoless* alternatives.

807

808

809

810

811

The response that Rubin's letter generated was more interesting than the letter itself. For five months, *Communications of the ACM* published letters that offered different versions of Rubin's original seven-line program. The letters were evenly divided between those defending *gotos* and those castigating them. Readers suggested roughly 17 different rewrites, and the rewritten code fully covered the spectrum of approaches to avoiding *gotos*. The editor of *CACM* noted that the letter had generated more response by far than any other issue ever considered in the pages of *CACM*.

812

813

814

815

816

817

818

819

For the follow-up letters, see

820

Communications of the ACM 30, no. 5 (May 1987): 351–55.

- 821 *Communications of the ACM* 30, no. 6 (June 1987): 475–78.
- 822 *Communications of the ACM* 30, no. 7 (July 1987): 632–34.
- 823 *Communications of the ACM* 30, no. 8 (August 1987): 659–62.
- 824 *Communications of the ACM* 30, no. 12 (December 1987): 997, 1085.
- 825 CC2E.COM/1706 Clark, R. Lawrence, “A Linguistic Contribution of GOTO-less Programming,”
- 826 *Datamation*, December 1973. This classic paper humorously argues for
- 827 replacing the “go to” statement with the “come from” statement. It was also
- 828 reprinted in the April 1974 edition of *Communications of the ACM*.

829 CC2E.COM/1713

CHECKLIST: Unusual Control Structures

830 ***return***

- 831 ☐ Does each routine use *return* only when necessary?
- 832 ☐ Do *returns* enhance readability?

833 **Recursion**

- 834 ☐ Does the recursive routine include code to stop the recursion?
- 835 ☐ Does the routine use a safety counter to guarantee that the routine stops?
- 836 ☐ Is recursion limited to one routine?
- 837 ☐ Is the routine’s depth of recursion within the limits imposed by the size of
- 838 the program’s stack?
- 839 ☐ Is recursion the best way to implement the routine? Is it better than simple
- 840 iteration?

841 ***goto***

- 842 ☐ Are *gotos* used only as a last resort, and then only to make code more
- 843 readable and maintainable?
- 844 ☐ If a *goto* is used for the sake of efficiency, has the gain in efficiency been
- 845 measured and documented?
- 846 ☐ Are *gotos* limited to one label per routine?
- 847 ☐ Do all *gotos* go forward, not backward?
- 848 ☐ Are all *goto* labels used?
-

Key Points

- Multiple *returns* can enhance a routine's readability and maintainability, and they help prevent deeply nested logic. They should, nevertheless, be used carefully.
- Recursion provides elegant solutions to a small set of problems. Use it carefully, too.
- In a few cases, *gotos* are the best way to write code that's readable and maintainable. Such cases are rare. Use *gotos* only as a last resort.