

# 23

## Debugging

**Contents**  
23.1 Overview of Debugging Issues  
23.2 Finding a Defect  
23.3 Fixing a Defect  
23.4 Psychological Considerations in Debugging  
23.5 Debugging Tools—Obvious and Not-So-Obvious

**Related Topics**  
The software-quality landscape: Chapter 20  
  
Developer testing: Chapter 22  
  
Refactoring: Chapter 24

DEBUGGING IS THE PROCESS OF IDENTIFYING the root cause of an error and correcting it. It contrasts with testing, which is the process of detecting the error initially. On some projects, debugging occupies as much as 50 percent of the total development time. For many programmers, debugging is the hardest part of programming.


Debugging doesn't have to be the hardest part. If you follow the advice in this book, you'll have fewer errors to debug. Most of the defects you will have will be minor oversights and typos, easily found by looking at a source-code listing or stepping through the code in a debugger. For the remaining harder bugs, this chapter describes how to make debugging much easier than it usually is

### 23.1 Overview of Debugging Issues

The late Rear Admiral Grace Hopper, co-inventor of COBOL, always said that the word "bug" in software dated back to the first large-scale digital computer, the Mark I (IEEE 1992). Programmers traced a circuit malfunction to the presence of a large moth that had found its way into the computer, and from that time on, computer problems were blamed on "bugs." Outside software, the word


“bug” dates back at least to Thomas Edison, who is quoted as using it as early as 1878 (Tenner 1997).

The word “bug” is a cute word and conjures up images like this one:



**G23xx01**

The reality of software defects, however, is that bugs aren’t organisms that sneak into your code when you forget to spray it with pesticide. They are errors. A bug in software means that a programmer made a mistake. The result of the mistake isn’t like the cute picture shown above. It’s more likely a note like this one:



**G23xx02**

In this context, technical accuracy requires that mistakes in the code be called “errors,” “defects,” or “faults.”

**Role of Debugging in Software Quality**

Like testing, debugging isn’t a way to improve the quality of your software, per se; it’s a way to diagnose defects. Software quality must be built in from the start. The best way to build a quality product is to develop requirements carefully, design well, and use high-quality coding practices. Debugging is a last resort.

## Variations in Debugging Performance

Why talk about debugging? Doesn't everyone know how to debug?

KEY POINT

No, not everyone knows how to debug. Studies of experienced programmers have found roughly a 20-to-1 difference in the time it takes experienced programmers to find the same set of defects. Moreover, some programmers find more defects and make corrections more accurately. Here are the results of a classic study that examined how effectively professional programmers with at least four years of experience debugged a program with 12 defects:

	Fastest Three Programmers	Slowest Three Programmers
Average debug time (minutes)	5.0	14.1
Average number of defects not found	0.7	1.7
Average number of defects made correcting defects	3.0	7.7

Source: "Some Psychological Evidence on How People Debug Computer Programs" (Gould 1975).

HARD DATA

The three programmers who were best at debugging were able to find the defects in about one-third the time and inserted only about two-fifths as many new defects as the three who were the worst. The best programmer found all the defects and didn't insert any new defects in correcting them. The worst missed 4 of the 12 defects and inserted 11 new defects in correcting the 8 defects he found.

But, this study doesn't really tell the whole story. After the first round of debugging the fastest three programmers still have 3.7 defects left in their code, and the slowest still have 9.4 defects. Neither group is done debugging yet. I wondered what would happen if I applied the same find-and-bad-fix ratios to additional debugging cycles. This isn't statistically valid, but it's still interesting. When I applied the same find-and-fix ratios to successive debugging cycles until each group had less than half a defect remaining, the fastest group required a total of 3 debugging cycles, whereas the slowest group required 14 debugging cycles. Bearing in mind that each cycle of the slower group takes almost 3 times as long as each cycle of the fastest group, the slowest group would take about 13 times as long to fully debug its programs as the fastest group, according to my non-scientific extrapolation of this study. Interestingly, this wide variation has been confirmed by other studies (Gilb 1977, Curtis 1981).

77 **CROSS-REFERENCE** For  
78 details on the relationship  
79 between quality and cost, see  
80 Section 20.5, “The General  
81 Principle of Software  
Quality.”

In addition to providing insight into debugging, the evidence supports the General Principle of Software Quality: Improving quality reduces development costs. The best programmers found the most defects, found the defects most quickly, and made correct modifications most often. You don’t have to choose between quality, cost, and time—they all go hand in hand.

## Defects as Opportunities

What does having an defect mean? Assuming that you don’t want the program to have defect, it means that you don’t fully understand what the program does. The idea of not understanding what the program does is unsettling. After all, if you created the program, it should do your bidding. If you don’t know exactly what you’re telling the computer to do, that’s only a small step from merely trying different things until something seems to work—that is, programming by trial and error. If you’re programming by trial and error, defects are guaranteed. You don’t need to learn how to fix defects; you need to learn how to avoid them in the first place.

Most people are somewhat fallible, however, and you might be an excellent programmer who has simply made a modest oversight. If this is the case, an error in your program represents a powerful opportunity. You can:

### *Learn about the program you’re working on*

You have something to learn about the program because if you already knew it perfectly, it wouldn’t have a defect. You would have corrected it already.

### *Learn about the kind of mistakes you make*

If you wrote the program, you inserted the defect. It’s not every day that a spotlight exposes a weakness with glaring clarity, but this particular day you have an opportunity to learn about your mistakes. Once you find the mistake, ask why did you make it? How could you have found it more quickly? How could you have prevented it? Does the code have other mistakes just like it? Can you correct them before they cause problems of their own?

### *Learn about the quality of your code from the point of view of someone who has to read it*

You’ll have to read your code to find the defect. This is an opportunity to look critically at the quality of your code. Is it easy to read? How could it be better? Use your discoveries to refactor your current code or to improve the code you write next.

### *Learn about how you solve problems*

Does your approach to solving debugging problems give you confidence? Does your approach work? Do you find defects quickly? Or is your approach to debugging weak? Do you feel anguish and frustration? Do you guess randomly?

99 **FURTHER READING** For  
100 details on practices that will  
101 help you learn about the  
102 kinds of errors you are  
103 personally prone to, see *A  
104 Discipline for Software  
Engineering* (Humphrey  
1995).

115 Do you need to improve? Considering the amount of time many projects spend  
116 on debugging, you definitely won't waste time if you observe how you debug.  
117 Taking time to analyze and change the way you debug might be the quickest way  
118 to decrease the total amount of time it takes you to develop a program.

119 ***Learn about how you fix defects***

120 In addition to learning how you find defects, you can learn about how you fix  
121 them. Do you make the easiest possible correction, by applying *goto* Band-Aids  
122 and special-case makeup that changes the symptom but not the problem? Or do  
123 you make systemic corrections, demanding an accurate diagnosis and prescribing  
124 treatment for the heart of the problem?

125 All things considered, debugging is an extraordinarily rich soil in which to plant  
126 the seeds of your own improvement. It's where all construction roads cross:  
127 readability, design, code quality—you name it. This is where building good code  
128 pays off—especially if you do it well enough that you don't have to debug very  
129 often.

130 **An Ineffective Approach**

131 Unfortunately, programming classes in colleges and universities hardly ever  
132 offer instruction in debugging. If you studied programming in college, you might  
133 have had a lecture devoted to debugging. Although my computer-science  
134 education was excellent, the extent of the debugging advice I received was to  
135 "put print statements in the program to find the defect." This is not adequate. If  
136 other programmers' educational experiences are like mine, a great many  
137 programmers are being forced to reinvent debugging concepts on their own.  
138 What a waste!

139 **The Devil's Guide to Debugging**

140 ***Programmers do not***  
141 ***always use available data***  
142 ***to constrain their***  
143 ***reasoning. They carry out***  
144 ***minor and irrational***  
145 ***repairs, and they often***  
146 ***don't undo the incorrect***  
147 ***repairs.***  
148 ***—Iris Vessey***

149 In Dante's vision of hell, the lowest circle is reserved for Satan himself. In  
150 modern times, Old Scratch has agreed to share the lowest circle with  
151 programmers who don't learn to debug effectively. He tortures programmers by  
making them use this common debugging approach:

***Find the defect by guessing***

To find the defect, scatter print statements randomly throughout a program.  
Examine the output to see where the defect is. If you can't find the defect with  
print statements, try changing things in the program until something seems to  
work. Don't back up the original version of the program, and don't keep a record  
of the changes you've made. Programming is more exciting when you're not  
quite sure what the program is doing. Stock up on Jolt cola and Twinkies  
because you're in for a long night in front of the terminal.

***Don't waste time trying to understand the problem***

It's likely that the problem is trivial, and you don't need to understand it completely to fix it. Simply finding it is enough.

***Fix the error with the most obvious fix***

It's usually good just to fix the specific problem you see, rather than wasting a lot of time making some big, ambitious correction that's going to affect the whole program. This is a perfect example:

```
x = Compute( y )
if ( y = 17 )
    x = $25.15      -- Compute() doesn't work for y = 17, so fix it
```

Who needs to dig all the way into *Compute()* for an obscure problem with the value of *17* when you can just write a special case for it in the obvious place?

This approach is infinitely extendable. If we later find that *Compute()* returns the wrong value when *y=18*, we just extend our fix:

```
x = Compute( y )
if ( y = 17 )
    x = $25.15      -- Compute() doesn't work for y = 17, so fix it
else if ( y = 18 )
    x = $27.85      -- Compute() doesn't work for y = 18, so fix it
```

**Debugging by Superstition**

Satan has leased part of hell to programmers who debug by superstition. Every group has one programmer who has endless problems with demon machines, mysterious compiler defects, hidden language defects that appear when the moon is full, bad data, losing important changes, a vindictive, possessed editor that saves programs incorrectly—you name it. This is “programming by superstition.”

If you have a problem with a program you've written, it's your fault. It's not the computer's fault, and it's not the compiler's fault. The program doesn't do something different every time. It didn't write itself; you wrote it, so take responsibility for it.

**KEY POINT**

Even if an error at first appears not to be your fault, it's strongly in your interest to assume that it is. That assumption helps you debug: It's hard enough to find a defect in your code when you're looking for it; it's even harder when you've assumed your code is error-free. It improves your credibility because when you do claim that an error arose from someone else's code, other programmers will believe that you have checked out the problem carefully. Assuming the error is your fault also saves you the embarrassment of claiming that an error is someone

189 else's fault and then having to recant publicly later when you find out that it was  
190 your defect after all.

## 191 23.2 Finding a Defect

192 Debugging consists of finding the defect and fixing it. Finding the defect (and  
193 understanding it) is usually 90 percent of the work.

194 Fortunately, you don't have to make a pact with Satan in order to find an  
195 approach to debugging that's better than random guessing. Contrary to what the  
196 Devil wants you to believe, debugging by thinking about the problem is much  
197 more effective and interesting than debugging with an eye of newt and the dust  
198 of a frog's ear.

199 Suppose you were asked to solve a murder mystery. Which would be more  
200 interesting: going door to door throughout the county, checking every person's  
201 alibi for the night of October 17, or finding a few clues and deducing the  
202 murderer's identity? Most people would rather deduce the person's identity, and  
203 most programmers find the intellectual approach to debugging more satisfying.  
204 Even better, the effective programmers who debug in one-twentieth the time of  
205 the ineffective programmers aren't randomly guessing about how to fix the  
206 program. They're using the scientific method.

## 207 The Scientific Method of Debugging

208 Here are the steps you go through when you use the scientific method:

- 209 1. Gather data through repeatable experiments.
- 210 2. Form a hypothesis that accounts for the relevant data.
- 211 3. Design an experiment to prove or disprove the hypothesis.
- 212 4. Prove or disprove the hypothesis.
- 213 5. Repeat as needed.

### 214 KEY POINT

215 This process has many parallels in debugging. Here's an effective approach for  
finding a defect:

- 216 1. Stabilize the error.
- 217 2. Locate the source of the error (the "fault").

- 218                   a. Gather the data that produces the defect.
- 219                   b. Analyze the data that has been gathered and form a hypothesis about the
- 220                   defect.
- 221                   c. Determine how to prove or disprove the hypothesis, either by testing the
- 222                   program or by examining the code.
- 223                   d. Prove or disprove the hypothesis using the procedure identified in 2(c).
- 224                   3. Fix the defect.
- 225                   4. Test the fix.
- 226                   5. Look for similar errors.

227                   The first step is similar to the scientific method's first step in that it relies on

228                   repeatability. The defect is easier to diagnose if you can make it occur reliably.

229                   The second step uses the first four steps of the scientific method. You gather the

230                   test data that divulged the defect, analyze the data that has been produced, and

231                   form a hypothesis about the source of the error. You design a test case or an

232                   inspection to evaluate the hypothesis and then declare success or renew your

233                   efforts, as appropriate.

234                   Let's look at each of the steps in conjunction with an example.

235                   Assume that you have an employee database program that has an intermittent

236                   error. The program is supposed to print a list of employees and their income-tax

237                   withholdings in alphabetical order. Here's part of the output:

238	Formatting, Fred Freeform	\$5,877
239	Goto, Gary	\$1,666
240	Modula, Mildred	\$10,788
241	Many-Loop, Mavis	\$8,889
242	Statement, Sue Switch	\$4,000
243	Whileloop, Wendy	\$7,860

244                   The error is that *Many-Loop*, *Mavis* and *Modula*, *Mildred* are out of order.

## 245                   Stabilize the Error

246                   If a defect doesn't occur reliably, it's almost impossible to diagnose. Making an

247                   intermittent defect occur predictably is one of the most challenging tasks in

248                   debugging.



249 **CROSS-REFERENCE** For  
250 details on using pointers  
251 safely, see Section 13.2,  
252 “Pointers.”

An error that doesn’t occur predictably is usually an initialization error or a dangling-pointer problem. If the calculation of a sum is right sometimes and wrong sometimes, a variable involved in the calculation probably isn’t being initialized properly—most of the time it just happens to start at 0. If the problem is a strange and unpredictable phenomenon and you’re using pointers, you almost certainly have an uninitialized pointer or are using a pointer after the memory that it points to has been deallocated.

Stabilizing an error usually requires more than finding a test case that produces the error. It includes narrowing the test case to the simplest one that still produces the error. If you work in an organization that has an independent test team, sometimes it’s the team’s job to make the test cases simple. Most of the time, it’s your job.

To simplify the test case, you bring the scientific method into play again. Suppose you have 10 factors that, used in combination, produce the error. Form a hypothesis about which factors were irrelevant to producing the error. Change the supposedly irrelevant factors, and rerun the test case. If you still get the error, you can eliminate those factors and you’ve simplified the test. Then you can try to simplify the test further. If you don’t get the error, you’ve disproved that specific hypothesis, and you know more than you did before. It might be that some subtly different change would still produce the error, but you know at least one specific change that does not.

In the employee withholdings example, when the program is run initially, *Many-Loop*, *Mavis* is listed after *Modula*, *Mildred*. When the program is run a second time, however, the list is fine:

Formatting, Fred Freeform	\$5,877
Goto, Gary	\$1,666
Many-Loop, Mavis	\$8,889
Modula, Mildred	\$10,788
Statement, Sue Switch	\$4,000
Whileloop, Wendy	\$7,860

It isn’t until *Fruit-Loop*, *Frita* is entered and shows up in an incorrect position that you remember that *Modula*, *Mildred* had been entered just before she showed up in the wrong spot too. What’s odd about both cases is that they were entered singly. Usually, employees are entered in groups.

You hypothesize: The problem has something to do with entering a single new employee.

If this is true, running the program again should put *Fruit-Loop*, *Frita* in the right position. Here’s the result of a second run:

287	Formatting, Fred Freeform	\$5,877
288	Fruit-Loop, Frita	\$5,771
289	Goto, Gary	\$1,666
290	Many-Loop, Mavis	\$8,889
291	Modula, Mildred	\$10,788
292	Statement, Sue Switch	\$4,000
293	Whileloop, Wendy	\$7,860

294 This successful run supports the hypothesis. To confirm it, you want to try  
 295 adding a few new employees, one at a time, to see whether they show up in the  
 296 wrong order and whether the order changes on the second run.

## 297 **Locate the Source of the Error**

298 The goal of simplifying the test case is to make it so simple that changing any  
 299 aspect of it changes the behavior of the error. Then, by changing the test case  
 300 carefully and watching the program's behavior under controlled conditions, you  
 301 can diagnose the problem.

302 Locating the source of the error also calls for using the scientific method. You  
 303 might suspect that the defect is a result of a specific problem, say an off-by-one  
 304 error. You could then vary the parameter you suspect is causing the problem—  
 305 one below the boundary, on the boundary, and one above the boundary—and  
 306 determine whether your hypothesis is correct.

307 In the running example, the source of the problem could be an off-by-one defect  
 308 that occurs when you add one new employee but not when you add two or more.  
 309 Examining the code, you don't find an obvious off-by-one defect. Resorting to  
 310 Plan B, you run a test case with a single new employee to see whether that's the  
 311 problem. You add *Hardcase, Henry* as a single employee and hypothesize that  
 312 his record will be out of order. Here's what you find:

313	Formatting, Fred Freeform	\$5,877
314	Fruit-Loop, Frita	\$5,771
315	Goto, Gary	\$1,666
316	Hardcase, Henry	\$493
317	Many-Loop, Mavis	\$8,889
318	Modula, Mildred	\$10,788
319	Statement, Sue Switch	\$4,000
320	Whileloop, Wendy	\$7,860

321 The line for *Hardcase, Henry* is exactly where it should be, which means that  
 322 your first hypothesis is false. The problem isn't caused simply by adding one  
 323 employee at a time. It's either a more complicated problem or something  
 324 completely different.

325 Examining the test-run output again, you notice that *Fruit-Loop, Frita* and  
 326 *Many-Loop, Mavis* are the only names containing hyphens. *Fruit-Loop* was out  
 327 of order when she was first entered, but *Many-Loop* wasn't, was she? Although

328 you don't have a printout from the original entry, in the original error *Modula*,  
329 *Mildred* appeared to be out of order, but she was next to *Many-Loop*. Maybe  
330 *Many-Loop* was out of order and *Modula* was all right.

331 You hypothesize: The problem arises from names with hyphens, not names that  
332 are entered singly.

333 But how does that account for the fact that the problem shows up only the first  
334 time an employee is entered? You look at the code and find that two different  
335 sorting routines are used. One is used when an employee is entered, and another  
336 is used when the data is saved. A closer look at the routine used when an  
337 employee is first entered shows that it isn't supposed to sort the data completely.  
338 It only puts the data in approximate order to speed up the save routine's sorting.  
339 Thus, the problem is that the data is printed before it's sorted. The problem with  
340 hyphenated names arises because the rough-sort routine doesn't handle niceties  
341 such as punctuation characters. Now, you can refine the hypothesis even further.

342 You hypothesize: Names with punctuation characters aren't sorted correctly until  
343 they're saved.

344 You later confirm this hypothesis with additional test cases.

## 345 Tips for Finding Defects

346 Once you've stabilized an error and refined the test case that produces it, finding  
347 its source can be either trivial or challenging, depending on how well you've  
348 written your code. If you're having a hard time finding a defect, it could be  
349 because the code isn't well written. You might not want to hear that, but it's true.  
350 If you're having trouble, consider these tips:

### 351 *Use all the data available to make your hypothesis*

352 When creating a hypothesis about the source of a defect, account for as much of  
353 the data as you can in your hypothesis. In the example, you might have noticed  
354 that *Fruit-Loop*, *Frita* was out of order and created a hypothesis that names  
355 beginning with an "F" are sorted incorrectly. That's a poor hypothesis because it  
356 doesn't account for the fact that *Modula*, *Mildred* was out of order or that names  
357 are sorted correctly the second time around. If the data doesn't fit the hypothesis,  
358 don't discard the data—ask why it doesn't fit, and create a new hypothesis.

359 The second hypothesis in the example, that the problem arises from names with  
360 hyphens, not names that are entered singly, didn't seem initially to account for  
361 the fact that names were sorted correctly the second time around either. In this  
362 case, however, the second hypothesis led to a more refined hypothesis that  
363 proved to be correct. It's all right that the hypothesis doesn't account for all of

364 the data at first as long as you keep refining the hypothesis so that it does  
365 eventually.

366 ***Refine the test cases that produce the error***

367 If you can't find the source of an error, try to refine the test cases further than  
368 you already have. You might be able to vary one parameter more than you had  
369 assumed, and focusing on one of the parameters might provide the crucial  
370 breakthrough.

371 **CROSS-REFERENCE** For  
372 more on unit test  
373 frameworks, see "Plug unit  
374 tests into a test framework"  
in Section 22.4.

375 ***Exercise the code in your unit test suite***

376 Defects tend to be easier to find in small fragments of code than in large  
377 integrated programs. Use your unit tests to test the code in isolation.

378 ***Use available tools***

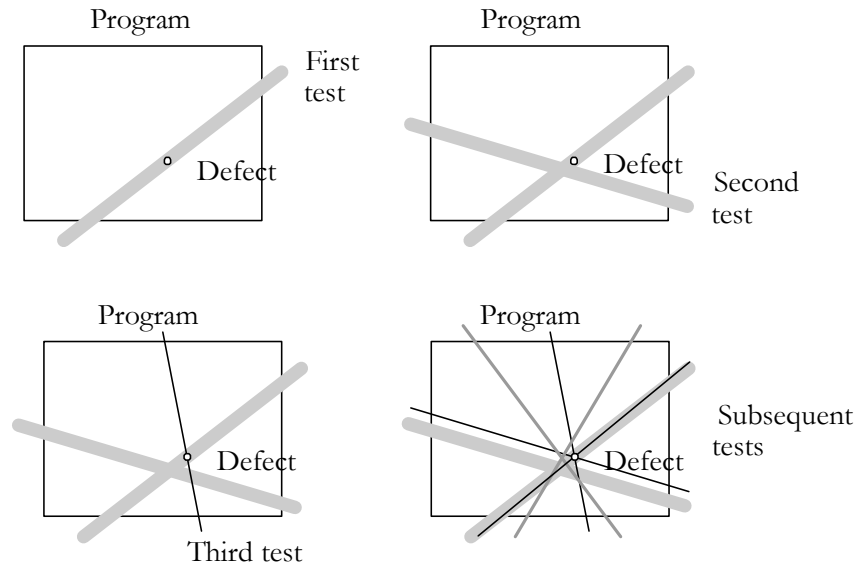
379 Numerous tools are available to support debugging sessions: interactive  
380 debuggers, picky compilers, memory checkers, and so on. The right tool can  
381 make a difficult job easy. With one tough-to-find error, for example, one part of  
382 the program was overwriting another part's memory. This error was difficult to  
383 diagnose using conventional debugging practices because the programmer  
couldn't determine the specific point at which the program was incorrectly  
overwriting memory. The programmer used a memory breakpoint to set a watch  
on a specific memory address. When the program wrote to that memory location,  
the debugger stopped the code, and the guilty code was exposed.

384 This is an example of problem that's difficult to diagnose analytically but which  
385 becomes quite simple when the right tool is applied.

386 ***Reproduce the error several different ways***

387 Sometimes trying cases that are similar to the error-producing case, but not  
388 exactly the same, is instructive. Think of this approach as triangulating the  
389 defect. If you can get a fix on it from one point and a fix on it from another, you  
390 can determine exactly where it is.

391 Reproducing the error several different ways helps diagnose the cause of the  
392 error. Once you think you've identified the defect, run a case that's close to the  
393 cases that produce errors but that should not produce an error itself. If it does  
394 produce an error, you don't completely understand the problem yet. Errors often  
395 arise from combinations of factors, and trying to diagnose the problem with only  
396 one test case sometimes doesn't diagnose the root problem.



### F23xx01

**Figure 23-1**

*Try to reproduce an error several different ways to determine its exact cause.*

#### ***Generate more data to generate more hypotheses***

Choose test cases that are different from the test cases you already know to be erroneous or correct. Run them to generate more data, and use the new data to add to your list of possible hypotheses.

#### ***Use the results of negative tests***

Suppose you create a hypothesis and run a test case to prove it. Suppose the test case disproves the hypothesis, so that you still don't know the source of the error. You still know something you didn't before—namely, that the defect is not in the area in which you thought it was. That narrows your search field and the set of possible hypotheses.

#### ***Brainstorm for possible hypotheses***

Rather than limiting yourself to the first hypothesis you think of, try to come up with several. Don't analyze them at first—just come up with as many as you can in a few minutes. Then look at each hypothesis and think about test cases that would prove or disprove it. This mental exercise is helpful in breaking the debugging logjam that results from concentrating too hard on a single line of reasoning.

#### ***Narrow the suspicious region of the code***

If you've been testing the whole program, or a whole class or routine, test a smaller part instead. Use print statements, logging, or tracing to identify which section of code is producing the error.

422 If you need a more powerful technique to narrow the suspicious region of the  
423 code, systematically remove parts of the program and see whether the error still  
424 occurs. If it doesn't, you know it's in the part you took away. If it does, you  
425 know it's in the part you've kept.

426 Rather than removing regions haphazardly, divide and conquer. Use a binary  
427 search algorithm to focus your search. Try to remove about half the code the first  
428 time. Determine the half the defect is in, and then divide that section. Again,  
429 determine which half contains the defect, and again, chop that section in half.  
430 Continue until you find the defect.

431 If you use many small routines, you'll be able to chop out sections of code  
432 simply by commenting out calls to the routines. Otherwise, you can use  
433 comments or preprocessor commands to remove code.

434 If you're using a debugger, you don't necessarily have to remove pieces of code.  
435 You can set a breakpoint partway through the program and check for the defect  
436 that way instead. If your debugger allows you to skip calls to routines, eliminate  
437 suspects by skipping the execution of certain routines and seeing whether the  
438 error still occurs. The process with a debugger is otherwise similar to the one in  
439 which pieces of a program are physically removed.

440 **CROSS-REFERENCE** For  
441 more details on error-prone  
442 code, see "Target error-prone  
443 modules" in Section 24.6.  
444

#### ***Be suspicious of classes and routines that have had defects before***

Classes that have had defects before are likely to continue to have defects. A class that has been troublesome in the past is more likely to contain a new defect than a class that has been defect-free. Re-examine error-prone classes and routines.

#### ***Check code that's changed recently***

445 If you have a new error that's hard to diagnose, it's usually related to code that's  
446 changed recently. It could be in completely new code or in changes to old code.  
447 If you can't find a defect, run an old version of the program to see whether the  
448 error occurs. If it doesn't, you know the error's in the new version or is caused  
449 by an interaction with the new version. Scrutinize the differences between the  
450 old and new versions. Check the version control log to see what code has  
451 changed recently. If that's not possible, use a diff tool to compare changes in the  
452 old, working source code to the new, broken source code.  
453

#### ***Expand the suspicious region of the code***

454 It's easy to focus on a small section of code, sure that "the defect *must* be in this  
455 section." If you don't find it in the section, consider the possibility that the defect  
456 isn't in the section. Expand the area of code you suspect, and then focus on  
457 pieces of it using the binary search technique described above.  
458

459 **CROSS-REFERENCE** For  
460 a full discussion of  
461 integration, see Chapter 29,  
“Integration.”

462  
463  
464  
465  
466  
467

### *Integrate incrementally*

Debugging is easy if you add pieces to a system one at a time. If you add a piece to a system and encounter a new error, remove the piece and test it separately.

### *Check for common defects*

Use code-quality checklists to stimulate your thinking about possible defects. If you’re following the inspection practices described in Section 21.3, you’ll have your own fine-tuned checklist of the common problems in your environment. You can also use the checklists that appear throughout this book. See the “List of Checklists” following the table of contents.

468 **CROSS-REFERENCE** For  
469 details on how involving  
470 other developers can put a  
471 beneficial distance between  
you and the problem, see  
Section 21.1, “Overview of  
472 Collaborative Development  
473 Practices.”

474  
475  
476  
477  
478  
479

### *Talk to someone else about the problem*

Some people call this “confessional debugging.” You often discover your own defect in the act of explaining it to another person. For example, if you were explaining the problem in the salary example, you might sound like this:

“Hey, Jennifer, have you got a minute? I’m having a problem. I’ve got this list of employee salaries that’s supposed to be sorted, but some names are out of order. They’re sorted all right the second time I print them out but not the first. I checked to see if it was new names, but I tried some that worked. I know they should be sorted the first time I print them because the program sorts all the names as they’re entered and again when they’re saved—wait a minute—no, it doesn’t sort them when they’re entered. That’s right. It only orders them roughly. Thanks, Jennifer. You’ve been a big help.”

Jennifer didn’t say a word, and you solved your problem. This result is typical, and this approach is perhaps your most potent tool for solving difficult defects.

### *Take a break from the problem*

Sometimes you concentrate so hard you can’t think. How many times have you paused for a cup of coffee and figured out the problem on your way to the coffee machine? Or in the middle of lunch? Or on the way home? Or in the shower the next morning? If you’re debugging and making no progress, once you’ve tried all the options, let it rest. Go for a walk. Work on something else. Go home for the day. Let your subconscious mind tease a solution out of the problem.

The auxiliary benefit of giving up temporarily is that it reduces the anxiety associated with debugging. The onset of anxiety is a clear sign that it’s time to take a break.

## **Brute Force Debugging**

Brute force is an often-overlooked approach to debugging software problems. By “brute force,” I’m referring to a technique that might be tedious, arduous, and time-consuming, but that it is *guaranteed* to solve the problem. Which specific

496 techniques are guaranteed to solve a problem are context dependent, but here are  
497 some general candidates:

- 498 ● Perform a full design and/or code review on the broken code
- 499 ● Throw away the section of code and redesign/recode it from scratch
- 500 ● Throw away the whole program and redesign/recode it from scratch
- 501 ● Compile code with full debugging information
- 502 ● Compile code at pickiest warning level and fix all the picky compiler  
503 warnings
- 504 ● Strap on a unit test harness and test the new code in isolation
- 505 ● Create an automated test suite and run it all night
- 506 ● Step through a big loop in the debugger manually until you get to the error  
507 condition
- 508 ● Instrument the code with print, display, or other logging statements
- 509 ● Replicate the end-user's full machine configuration
- 510 ● Integrate new code in small pieces, fully testing each piece as its integrated

#### 511 ***Set a maximum time for quick and dirty debugging***

512 For each brute force technique, your reaction might very well be, "I can't do  
513 that; it's too much work!" The point is that it's only too much work if it takes  
514 more time than what I call "quick and dirty debugging." It's always tempting to  
515 try for a quick guess rather than systematically instrumenting the code and  
516 giving the defect no place to hide. The gambler in each of us would rather use a  
517 risky approach that might find the defect in five minutes than the surefire  
518 approach that will find the defect in half an hour. The risk is that, if the five-  
519 minute approach doesn't work, you get stubborn. Finding the defect the "easy"  
520 way becomes a matter of principle, and hours pass unproductively, as do days,  
521 weeks, months, ... How often have you spent two hours debugging code that took  
522 only 30 minutes to write? That's a bad distribution of labor, and you would have  
523 been better off simply to rewrite the code than to debug bad code.

524 When you decide to go for the quick victory, set a maximum time limit for trying  
525 the quick way. If you go past the time limit, resign yourself to the idea that the  
526 defect is going to be harder to diagnose than you originally thought, and flush it  
527 out the hard way. This approach allows you to get the easy defects right away  
528 and the hard defects after a bit longer.

#### 529 ***Make a list of brute force techniques***

530 Before you begin debugging a difficult error, ask yourself, "If I get stuck  
531 debugging this problem, is there some way that I am *guaranteed* to be able to fix



the problem?” If you can identify at least one brute force technique that will fix the problem—including rewriting the code in question—it’s less likely that you’ll waste hours or days when there’s a quicker alternative.

## Syntax Errors

Syntax-error problems are going the way of the woolly mammoth and the saber-toothed tiger. Compilers are getting better at diagnostic messages, and the days when you had to spend two hours finding a misplaced semicolon in a Pascal listing are almost gone. Here’s a list of guidelines you can use to hasten the extinction of this endangered species:

### *Don’t trust line numbers in compiler messages*

When your compiler reports a mysterious syntax error, look immediately before and immediately after the error—the compiler could have misunderstood the problem or simply have poor diagnostics. Once you find the real defect, try to determine the reason the compiler put the message on the wrong statement. Understanding your compiler better can help you find future defects.

### *Don’t trust compiler messages*

Compilers try to tell you exactly what’s wrong, but compilers are dissembling little rascals, and you often have to read between the lines to know what one really means. For example, in UNIX C, you can get a message that says “floating exception” for an integer divide-by-0. With C++’s Standard Template Library, you can get a pair of error messages: the first message is the real error in the use of the STL; the second message is a message from the compiler saying, “Error message too long for printer to print; message truncated.” You can probably come up with many examples of your own.

### *Don’t trust the compiler’s second message*

Some compilers are better than others at detecting multiple errors. Some compilers get so excited after detecting the first error that they become giddy and overconfident; they prattle on with dozens of error messages that don’t mean anything. Other compilers are more levelheaded, and although they must feel a sense of accomplishment when they detect an error, they refrain from spewing out inaccurate messages. If you can’t quickly find the source of the second or third error message, don’t worry about it. Fix the first one and recompile.

### *Divide and conquer*

The idea of dividing the program into sections to help detect defects works especially well for syntax errors. If you have a troublesome syntax error, remove part of the code and compile again. You’ll either get no error (because the error’s in the part you removed), get the same error (meaning you need to remove a

different part), or get a different error (because you'll have tricked the compiler into producing a message that makes more sense).

**CROSS-REFERENCE** Many programming text editors can automatically find matching braces or *begin-end* pairs. For details on programming editors, see “Editing” in Section 30.2.

**Find extra comments and quotation marks**  
If your code is tripping up the compiler because it contains an extra quotation mark or beginning comment somewhere, insert the following sequence systematically into your code to help locate the defect:

```
C/C++/Java      /*"/**/
```

### 23.3 Fixing a Defect

The hard part is finding the defect. Fixing the defect is the easy part. But as with many easy tasks, the fact that it's easy makes it especially error-prone. At least one study found that defect corrections have more than a 50 percent chance of being wrong the first time (Yourdon 1986b). Here are a few guidelines for reducing the chance of error:

**KEY POINT**

**Understand the problem before you fix it**  
“The Devil’s Guide to Debugging” is right: The best way to make your life difficult and corrode the quality of your program is to fix problems without really understanding them. Before you fix a problem, make sure you understand it to the core. Triangulate the defect both with cases that should reproduce the error and with cases that shouldn’t reproduce the error. Keep at it until you understand the problem well enough to predict its occurrence correctly every time.

**HARD DATA**

**Understand the program, not just the problem**  
If you understand the context in which a problem occurs, you’re more likely to solve the problem completely rather than only one aspect of it. A study done with short programs found that programmers who achieve a global understanding of program behavior have a better chance of modifying it successfully than programmers who focus on local behavior, learning about the program only as they need to (Littman et al. 1986). Because the program in this study was small (280 lines), it doesn’t prove that you should try to understand a 50,000-line program completely before you fix a defect. It does suggest that you should understand at least the code in the vicinity of the defect correction—the “vicinity” being not a few lines but a few hundred.

**Confirm the defect diagnosis**  
Before you rush to fix a defect, make sure that you’ve diagnosed the problem correctly. Take the time to run test cases that prove your hypothesis and disprove competing hypotheses. If you’ve proven only that the error could be the result of

605 one of several causes, you don't yet have enough evidence to work on the one  
606 cause; rule out the others first.

607 ***Relax***

608 A programmer was ready for a ski trip. His product was ready to ship, he was  
609 already late, and he had only one more defect to correct. He changed the source  
610 file and checked it into version control. He didn't recompile the program and  
611 didn't verify that the change was correct.

612 ***Never debug standing up.*** In fact, the change was not correct, and his manager was outraged. How could he  
613 ***—Gerald Weinberg*** change code in a product that was ready to ship without checking it? What could  
614 be worse? Isn't this the pinnacle of professional recklessness?

615 If this isn't the height of recklessness, it's close, and it's common. Hurrying to  
616 solve a problem is one of the most time-ineffective things you can do. It leads to  
617 rushed judgments, incomplete defect diagnosis, and incomplete corrections.  
618 Wishful thinking can lead you to see solutions where there are none. The  
619 pressure—often self-imposed—encourages haphazard trial-and-error solutions,  
620 sometimes assuming that a solution works without verifying that it does.

621 In striking contrast, during the final days of Microsoft Windows 2000  
622 development, a developer needed to fix a defect that was the last remaining  
623 defect before a Release Candidate could be created. The developer changed the  
624 code, checked his fix, and tested his fix on his local build. But he didn't check  
625 the fix into version control at that point. Instead, he went to play basketball. He  
626 said, "I'm feeling too stressed right now to be sure that I've considered  
627 everything I should consider. I'm going to clear my mind for an hour, and then  
628 I'll come back and check in the code—once I've convinced myself that the fix is  
629 really correct."

630 Relax long enough to make sure your solution is right. Don't be tempted to take  
631 shortcuts. It might take more time, but it'll probably take less. If nothing else,  
632 you'll fix the problem correctly and your manager won't call you back from your  
633 ski trip.

634 **CROSS-REFERENCE** Gen  
635 eral issues involved in  
636 changing code are discussed  
637 in depth in Chapter 24,  
638 "Refactoring."

634 ***Save the original source code***  
635 Before you begin fixing the defect, be sure to archive a version of the code that  
636 you can return to later. It's easy to forget which change in a group of changes is  
637 the significant one. If you have the original source code, at least you can  
638 compare the old and the new files and see where the changes are.

639 ***Fix the problem, not the symptom***

640 You should fix the symptom too, but the focus should be on fixing the  
641 underlying problem rather than wrapping it in programming duct tape. If you

642  
643

don't thoroughly understand the problem, you're not fixing the code. You're fixing the symptom and making the code worse. Suppose you have this code:

644

### Java Example of Code That Needs to Be Fixed

```
for ( claimNumber = 0; claimNumber < numClaims[ client ]; claimNumber++ ) {
    sum[ client ] = sum[ client ] + claimAmount[ claimNumber ];
}
```

645  
646  
647  
648  
649

Further suppose that when *client* equals 45, *sum* turns out to be wrong by \$3.45. Here's the wrong way to fix the problem:

### 650 CODING HORROR

### Java Example of Making the Code Worse by "Fixing" It

```
for ( claimNumber = 0; claimNumber < numClaims[ client ]; claimNumber++ ) {
    sum[ client ] = sum[ client ] + claimAmount[ claimNumber ];
}

if ( client == 45 ) {
    sum[ 45 ] = sum[ 45 ] + 3.45;
}
```

651  
652  
653  
654  
655  
656  
657  
658  
659

Here's the "fix."

Now suppose that when *client* equals 37 and the number of claims for the client is 0, you're not getting 0. Here's the wrong way to fix the problem:

### 660 CODING HORROR

### Java Example of Making the Code Worse by "Fixing" It (continued)

```
for ( claimNumber = 0; claimNumber < numClaims[ client ]; claimNumber++ ) {
    sum[ client ] = sum[ client ] + claimAmount[ claimNumber ];
}

if ( client == 45 ) {
    sum[ 45 ] = sum[ 45 ] + 3.45;
}

else if ( ( client == 37 ) && ( numClaims[ client ] == 0 ) ) {
    sum[ 37 ] = 0.0;
}
```

661  
662  
663  
664  
665  
666  
667  
668  
669  
670  
671  
672  
673  
674

Here's the second "fix."

If this doesn't send a cold chill down your spine, you won't be affected by anything else in this book either. It's impossible to list all the problems with this approach in a book that's only a little over 900 pages long, but here are the top three:

675  
676  
677  
678  
679

- The fixes won't work most of the time. The problems look as though they're the result of initialization defects. Initialization defects are, by definition, unpredictable, so the fact that the sum for client 45 is off by \$3.45 today doesn't tell you anything about tomorrow. It could be off by \$10,000.02, or it could be correct. That's the nature of initialization defects.

- It's unmaintainable. When code is special-cased to work around errors, the special cases become the code's most prominent feature. The \$3.45 won't always be \$3.45, and another error will show up later. The code will be modified again to handle the new special case, and the special case for \$3.45 won't be removed. The code will become increasingly barnacled with special cases. Eventually the barnacles will be too heavy for the code to support, and the code will sink to the bottom of the ocean— a fitting place for it.
- It uses the computer for something that's better done by hand. Computers are good at predictable, systematic calculations, but humans are better at fudging data creatively. You'd be wiser to treat the output with Whiteout and a typewriter than to monkey with the code.

### *Change the code only for good reason*

Related to fixing symptoms is the technique of changing code at random until it seems to work. The typical line of reasoning goes like this: "This loop seems to contain a defect. It's probably an off-by-one error, so I'll just put a *-1* here and try it. OK. That didn't work, so I'll just put a *+1* in instead. OK. That seems to work. I'll say it's fixed."

As popular as this practice is, it isn't effective. Making changes to code randomly is like poking a Pontiac Aztek with a stick to see if it moves. You're not learning anything; you're just goofing around. By changing the program randomly, you say in effect, "I don't know what's happening here, but I'll try this change and hope it works." Don't change code randomly. That's voodoo programming. The more different you make it without understanding it, the less confidence you'll have that it works correctly.

Before you make a change, be confident that it will work. Being wrong about a change should leave you astonished. It should cause self-doubt, personal reevaluation, and deep soul-searching. It should happen rarely.

### *Make one change at a time*

Changes are tricky enough when they're done one at a time. When done two at a time, they can introduce subtle errors that look like the original errors. Then you're in the awkward position of not knowing whether (1) you didn't correct the error, (2) you corrected the error but introduced a new one that looks similar, or (3) you didn't correct the error and you introduced a similar new error. Keep it simple: Make just one change at a time.

715 **CROSS-REFERENCE** For  
716 details on automated  
717 regression testing, see  
718 “Retesting (Regression  
719 Testing)” in Section 22.6.  
720

### *Check your fix*

Check the program yourself, have someone else check it for you, or walk through it with someone else. Run the same triangulation test cases you used to diagnose the problem to make sure that all aspects of the problem have been resolved. If you’ve solved only part of the problem, you’ll find out that you still have work to do.

Rerun the whole program to check for side effects of your changes. The easiest and most effective way to check for side effects is to run the program through an automated suite of regression tests in JUnit, CppUnit, or equivalent.

### *Look for similar defects*

When you find one defect, look for others that are similar. Defects tend to occur in groups, and one of the values of paying attention to the kinds of defects you make is that you can correct all the defects of that kind. Looking for similar defects requires you to have a thorough understanding of the problem. Watch for the warning sign: If you can’t figure out how to look for similar defects, that’s a sign that you don’t yet completely understand the problem.

## 23.4 Psychological Considerations in Debugging

733 **FURTHER READING** For an  
734 excellent discussion of  
735 psychological issues in  
736 debugging, as well as many  
737 other areas of software  
738 development, see *The*  
739 *Psychology of Computer*  
740 *Programming* (Weinberg  
741 1998).

Debugging is as intellectually demanding as any other software-development activity. Your ego tells you that your code is good and doesn’t have a defect even when you have seen that it has one. You have to think precisely—forming hypotheses, collecting data, analyzing hypotheses, and methodically rejecting them—with a formality that’s unnatural to many people. If you’re both building code and debugging it, you have to switch quickly between the fluid, creative thinking that goes with design and the rigidly critical thinking that goes with debugging. As you read your code, you have to battle the code’s familiarity and guard against seeing what you expect to see.

### How “Psychological Set” Contributes to Debugging Blindness

When you see a token in a program that says *Num*, what do you see? Do you see a misspelling of the word “Numb”? Or do you see the abbreviation for “Number”? Most likely, you see the abbreviation for “Number.” This is the phenomenon of “psychological set”—seeing what you expect to see. What does this sign say?



### G23xx03

In this classic puzzle, people often see only one “the.” People see what they expect to see. Consider the following:

- Students learning *while* loops often expect a loop to be continuously evaluated; that is, they expect the loop to terminate as soon as the *while* condition becomes false, rather than only at the top or bottom (Curtis et al. 1986). They expect a *while* loop to act as “while” does in natural language.
- A programmer who unintentionally used both the variable *SYSTSTS* and the variable *SYSSSTS* thought he was using a single variable. He didn’t discover the problem until the program had been run hundreds of times, and a book was written containing the erroneous results (Weinberg 1998).
- A programmer looking at code like this code:

```
if ( x < y )
    swap = x
    x = y
    y = swap
```

sometimes sees code like this code:

```
if ( x < y ) {
    swap = x
    x = y
    y swap
}
```

People expect a new phenomenon to resemble similar phenomena they’ve seen before. They expect a new control construct to work the same as old constructs; programming-language *while* statements to work the same as real-life “while” statements; and variable names to be the same as they’ve been before. You see what you expect to see and thus overlook differences, like the misspelling of the word “language” in the previous sentence.

What does psychological set have to do with debugging? First, it speaks to the importance of good programming practices. Good formatting, commenting, variable names, routine names, and other elements of programming style help structure the programming background so that likely defects appear as variations and stand out.

#### HARD DATA

783 The second impact of psychological set is in selecting parts of the program to  
784 examine when an error is found. Research has shown that the programmers who  
785 debug most effectively mentally slice away parts of the program that aren't  
786 relevant during debugging (Basili, Selby, and Hutchens 1986). In general, the  
787 practice allows excellent programmers to narrow their search fields and find  
788 defects more quickly. Sometimes, however, the part of the program that contains  
789 the defect is mistakenly sliced away. You spend time scouring a section of code  
790 for a defect, and you ignore the section that contains the defect.

791 You took a wrong turn at the fork in the road and need to back up before you can  
792 go forward again. Some of the suggestions in Section 23.2's discussion of tips  
793 for finding defects are designed to overcome this "debugging blindness."

794 **How "Psychological Distance" Can Help**

795 **CROSS-REFERENCE** For  
796 details on creating variable  
797 names that won't be  
798 confusing, see Section 11.7,  
799 "Kinds of Names to Avoid."  
800 Psychological distance can be defined as the ease with which two items can be  
801 differentiated. If you are looking at a long list of words and have been told that  
802 they're all about ducks, you could easily mistake "Queck" for "Quack" because  
the two words look similar. The psychological distance between the words is  
small. You would be much less likely to mistake "Tuack" for "Quack" even  
though the difference is only one letter again. "Tuack" is less like "Quack" than  
"Queck" is because the first letter in a word is more prominent than the one in  
the middle.

803 Here are examples of psychological distances between variable names:

804 **Table 23-1. Examples of Psychological Distance Between Variable**  
805 **Names**

First Variable	Second Variable	Psychological Distance
stoppt	stcppt	Almost invisible
shiftrn	shiftrm	Almost none
dcount	bcoun	Small
claims1	claims2	Small
product	sum	Large

806 As you debug, be ready for the problems caused by insufficient psychological  
807 distance between similar variable names and between similar routine names. As  
808 you construct code, choose names with large differences so that you avoid the  
809 problem.



## 23.5 Debugging Tools—Obvious and Not-So-Obvious

**CROSS-REFERENCE** The line between testing tools and debugging tools is fuzzy. For details on testing tools, see Section 22.5, “Test-Support Tools.” For details on tools for other software-development activities, see Chapter 30, “Programming Tools.”

You can do much of the detailed, brain-busting work of debugging with debugging tools that are readily available. The tool that will drive the final stake through the heart of the defect vampire isn’t yet available, but each year brings an incremental improvement in available capabilities.

### Diff

A source-code comparator such as Diff is useful when you’re modifying a program in response to errors. If you make several changes and need to remove some that you can’t quite remember, a comparator can pinpoint the differences and jog your memory. If you discover a defect in a new version that you don’t remember in an older version, you can compare the files to determine exactly what changed.

### Compiler Warning Messages

One of the simplest and most effective debugging tools is your own compiler.

#### *Set your compiler’s warning level to the highest, pickiest level possible and fix the code so that it doesn’t produce any compiler warnings*

It’s sloppy to ignore compiler errors. It’s even sloppier to turn off the warnings so that you can’t even see them. Children sometimes think that if they close their eyes and can’t see you, they’ve made you go away. Setting a switch on the compiler to turn off warnings just means you can’t see the errors. It doesn’t make them go away any more than closing your eyes makes an adult go away.

Assume that the people who wrote the compiler know a great deal more about your language than you do. If they’re warning you about something, it usually means you have an opportunity to learn something new about your language. Make the effort to understand what the warning really means.

#### *Treat warnings as errors*

Some compilers let you treat warnings as errors. One reason to use the feature is that it elevates the apparent importance of a warning. Just as setting your watch five minutes fast tricks you into thinking it’s five minutes later than it is, setting your compiler to treat warnings as errors tricks you into taking them more seriously. Another reason to treat warnings as errors is that they often affect how your program compiles. When you compile and link a program, warnings typically won’t stop the program from linking but errors typically will. If you want to check warnings before you link, set the compiler switch that treats warnings as errors.

### *Initiate project wide standards for compile-time settings*

Set a standard that requires everyone on your team to compile code using the same compiler settings. Otherwise, when you try to integrate code compiled by different people with different settings, you'll get a flood of error messages and an integration nightmare.

## Extended Syntax and Logic Checking

You can use additional tools to check your code more thoroughly than your compiler does. For example, for C programmers, the lint utility painstakingly checks for use of uninitialized variables, writing = when you mean ==, and similarly subtle problems.

## Execution Profiler

You might not think of an execution profiler as a debugging tool, but a few minutes spent studying a program profile can uncover some surprising (and hidden) defects.

For example, I had suspected that a memory-management routine in one of my programs was a performance bottleneck. Memory management had originally been a small component using a linearly ordered array of pointers to memory. I replaced the linearly ordered array with a hash table in the expectation that execution time would drop by at least half. But after profiling the code, I found no change in performance at all. I examined the code more closely and found a defect that was wasting a huge amount of time in the allocation algorithm. The bottleneck hadn't been the linear-search technique; it was the defect. I hadn't needed to optimize the search after all. Examine the output of an execution profiler to satisfy yourself that your program spends a reasonable amount of time in each area.

## Test Frameworks/Scaffolding

As mentioned in Section 23.2 on finding defects, pulling out a troublesome piece of code, writing code to test it, and executing it by itself is often the most effective way to exorcise the demons from an error-prone program.

## Debugger

Commercially available debuggers have advanced steadily over the years, and the capabilities available today can change the way you program.

Good debuggers allow you to set breakpoints to break when execution reaches a specific line, or the  $n$ th time it reaches a specific line, or when a global variable changes, or when a variable is assigned a specific value. They allow you to step

**CROSS-REFERENCE** For details on scaffolding, see "Building Scaffolding to Test Individual Classes" in Section 22.5.

881 through code line by line, stepping through or over routines. They allow the  
882 program to be executed backwards, stepping back to the point where a defect  
883 originated. They allow you to log the execution of specific statements—similar  
884 to scattering “I’m here!” print statements throughout a program.

885 Good debuggers allow full examination of data, including structured and  
886 dynamically allocated data. They make it easy to view the contents of a linked  
887 list of pointers or a dynamically allocated array. They’re intelligent about user-  
888 defined data types. They allow you to make ad hoc queries about data, assign  
889 new values, and continue program execution.

890 You can look at the high-level language or the assembly language generated by  
891 your compiler. If you’re using several languages, the debugger automatically  
892 displays the correct language for each section of code. You can look at a chain of  
893 calls to routines and quickly view the source code of any routine. You can  
894 change parameters to a program within the debugger environment.

895 The best of today’s debuggers also remember debugging parameters  
896 (breakpoints, variables being watched, and so on) for each individual program so  
897 that you don’t have to re-create them for each program you debug.

898 System debuggers operate at the systems level rather than the applications level  
899 so that they don’t interfere with the execution of the program being debugged.  
900 They’re essential when you are debugging programs that are sensitive to timing  
901 or the amount of memory available.

902 *An interactive debugger*  
903 *is an outstanding*  
904 *example of what is not*  
905 *needed—it encourages*  
906 *trial-and-error hacking*  
907 *rather than systematic*  
908 *design, and also hides*  
909 *marginal people barely*  
910 *qualified for precision*  
911 *programming.*  
912 *—Harlan Mills*

Given the enormous power offered by modern debuggers, you might be surprised that anyone would criticize them. But some of the most respected people in computer science recommend not using them. They recommend using your brain and avoiding debugging tools altogether. Their argument is that debugging tools are a crutch and that you find problems faster by thinking about them than by relying on tools. They argue that you, rather than the debugger, should mentally execute the program to flush out defects.

Regardless of the empirical evidence, the basic argument against debuggers isn’t valid. The fact that a tool can be misused doesn’t imply that it should be rejected. You wouldn’t avoid taking aspirin merely because it’s possible to overdose. You wouldn’t avoid mowing your lawn with a power mower just because it’s possible to cut yourself. Any other powerful tool can be used or abused, and so can a debugger.

#### 915 KEY POINT

The debugger isn’t a substitute for good thinking. But, in some cases, thinking isn’t a substitute for a good debugger either. The most effective combination is good thinking and a good debugger.

CC2E.COM/2368

918

---

**CHECKLIST: Debugging Reminders**

---

919

**Techniques for Finding Defects**

920

☐ Use all the data available to make your hypothesis

921

☐ Refine the test cases that produce the error

922

☐ Exercise the code in your unit test suite

923

☐ Use available tools

924

☐ Reproduce the error several different ways

925

☐ Generate more data to generate more hypotheses

926

☐ Use the results of negative tests

927

☐ Brainstorm for possible hypotheses

928

☐ Narrow the suspicious region of the code

929

☐ Be suspicious of classes and routines that have had defects before

930

☐ Check code that's changed recently

931

☐ Expand the suspicious region of the code

932

☐ Integrate incrementally

933

☐ Check for common defects

934

☐ Talk to someone else about the problem

935

☐ Take a break from the problem

936

☐ Set a maximum time for quick and dirty debugging

937

☐ Make a list of brute force techniques, and use them

938

**Techniques for Syntax Errors**

939

☐ Don't trust line numbers in compiler messages

940

☐ Don't trust compiler messages

941

☐ Don't trust the compiler's second message

942

☐ Divide and conquer

943

☐ Find extra comments and quotation marks

944

**Techniques for Fixing Defects**

945

☐ Understand the problem before you fix it

946

☐ Understand the program, not just the problem

947

☐ Confirm the defect diagnosis

948

☐ Relax

949

☐ Save the original source code

950

☐ Fix the problem, not the symptom

- 951 ☐ Change the code only for good reason
- 952 ☐ Make one change at a time
- 953 ☐ Check your fix
- 954 ☐ Look for similar defects

### 955 **General Approach to Debugging**

- 956 ☐ Do you use debugging as an opportunity to learn more about your program, mistakes, code quality, and problem-solving approach?
  - 957
  - 958 ☐ Do you avoid the trial-and-error, superstitious approach to debugging?
  - 959 ☐ Do you assume that errors are your fault?
  - 960 ☐ Do you use the scientific method to stabilize intermittent errors?
  - 961 ☐ Do you use the scientific method to find defects?
  - 962 ☐ Rather than using the same approach every time, do you use several different techniques to find defects?
  - 963
  - 964 ☐ Do you verify that the fix is correct?
  - 965 ☐ Do you use compiler warning messages, execution profiling, a test framework, scaffolding, and interactive debugging?
  - 966
  - 967
- 

CC2E.COM/2375

## 968 **Additional Resources**

969 Agans, David J. *Debugging: The Nine Indispensable Rules for Finding Even the*  
970 *Most Elusive Software and Hardware Problems*. Amacom, 2003. This book  
971 provides general debugging principles that can be applied in any language or  
972 environment.

973 Myers, Glenford J. *The Art of Software Testing*. New York: John Wiley, 1979.  
974 Chapter 7 of this classic book is devoted to debugging.

975 Allen, Eric. *Bug Patterns In Java*. Berkeley, Ca.: Apress, 2002. This book lays  
976 out an approach to debugging Java programs that is conceptually very similar to  
977 what is described in this chapter, including “The Scientific Method of  
978 Debugging,” distinguishing between debugging and testing, and identifying  
979 common bug patterns.

980 The following two books are similar in that their titles suggest they are  
981 applicable only to Microsoft Windows and .NET programs, but they both  
982 contain discussions of debugging in general, use of assertions, and coding  
983 practices that help to avoid bugs in the first place.

984 Robbins, John. *Debugging Applications for Microsoft .NET and Microsoft*  
985 *Windows*. Redmond, Wa.: Microsoft Press, 2003.

986 McKay, Everett N. and Mike Woodring, *Debugging Windows Programs:*  
987 *Strategies, Tools, and Techniques for Visual C++ Programmers*. Boston, Mass.:  
988 Addison Wesley, 2000.

## 989 Key Points

- 990 • Debugging is a make-or-break aspect of software development. The best  
991 approach is to use other techniques described in this book to avoid defects in  
992 the first place. It's still worth your time to improve your debugging skills,  
993 however, because the difference between good and poor debugging  
994 performance is at least 10 to 1.
- 995 • A systematic approach to finding and fixing errors is critical to success.  
996 Focus your debugging so that each test moves you a step forward. Use the  
997 Scientific Method of Debugging.
- 998 • Understand the root problem before you fix the program. Random guesses  
999 about the sources of errors and random corrections will leave the program in  
1000 worse condition than when you started.
- 1001 • Set your compiler warning to the pickiest level possible, and fix the errors it  
1002 reports. It's hard to fix subtle errors if you ignore the obvious ones.
- 1003 • Debugging tools are powerful aids to software development. Find them and  
1004 use them. Remember to use your brain at the same time.